



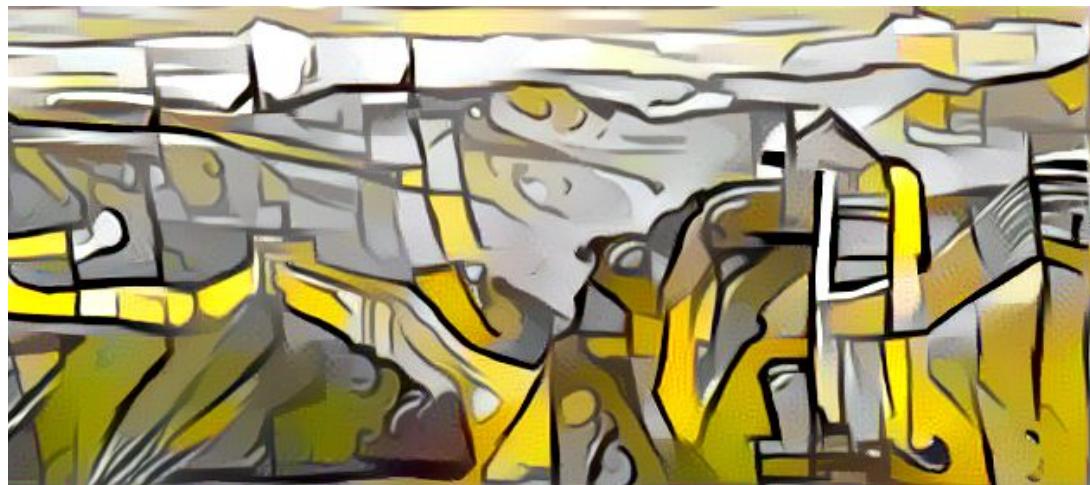
Getting up and running with TensorFlow

Create a **Kaggle.com** account

Sept. 18, 2017



Yufeng Guo
Developer Advocate,
Machine Learning
[@YufengG](https://twitter.com/YufengG)
yufengg.com





Yufeng
yfg@google.com
@YufengG

These slides:
bit.ly/tf-workshop-aicon

Break at 10:30am

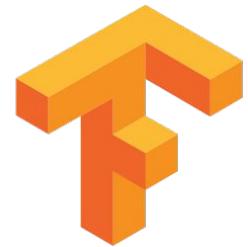
Create a Kaggle.com account



Overview of the Workshop

- Intro to TensorFlow, its ecosystem, cool stuff being done
 - (*& a little bit of keras*)
- TensorFlow's high-level 'canned' Estimators for Linear and Deep Neural Networks
 - *Features: Mushrooms and Fashion*
- “Wide & Deep” models with TensorFlow, and feature engineering
- Building Custom Estimators (for CNN models)

Create a Kaggle.com account



These slides: bit.ly/tf-workshop-aicon

Create a Kaggle.com account

Install TensorFlow & Jupyter: www.tensorflow.org/install/
pip install tensorflow pandas jupyter

Clone GitHub Repo: bit.ly/tensorflow-workshop

Optional: download the ‘Fashion-MNIST’ files:

<https://github.com/zalandoresearch/fashion-mnist#get-the-data>

TensorFlow Research, TensorFlow Fun



TensorFlow

tensorflow.org

- Fast, flexible, and scalable open-source machine learning library
- For research and production
- Apache 2.0 license

 tensorflow

Pull requests Issues Marketplace Gist

Repositories 13K Code 1M Commits 156K

13,398 repository results

tensorflow/tensorflow
Computation using data flow graphs for scalable machine learning
[tensorflow](#) [python](#) [machine-learning](#)
Updated an hour ago

rstudio/tensorflow
TensorFlow for R
Updated 14 hours ago

yao62995/tensorflow
图解tensorflow 源码
Updated on Nov 11, 2016

lmb-freiburg/hand3d
Network estimating 3D Handpose from single color images
[tensorflow](#)
Updated on May 4

mfigurnov/sact
Spatially Adaptive Computation Time for Residual Networks
Spatially Adaptive Computation Time for Residual Networks

 tensorflow

Pull requests Issues Marketplace Explore

Repositories 17K Code 2M Commits 206K Issues 33K Wikis 2K Users 240 Advanced search

17,890 repository results Sort: Best match ▾

tensorflow/tensorflow ● C++ ★ 70k
Computation using data flow graphs for scalable machine learning
[tensorflow](#) [python](#) [machine-learning](#)
Updated an hour ago

rstudio/tensorflow ● R ★ 853
TensorFlow for R
Updated 2 days ago

yao62995/tensorflow ★ 633
图解tensorflow 源码
Updated on Nov 11, 2016

lmb-freiburg/hand3d ● Python ★ 155
Network estimating 3D Handpose from single color images
[tensorflow](#)
Updated on Jul 4

mfigurnov/sact ● Python ★ 174
Spatially Adaptive Computation Time for Residual Networks
Spatially Adaptive Computation Time for Residual Networks

Languages

Python	9,600
Jupyter Notebook	3,563
HTML	405
C++	340
Shell	213
Java	172
JavaScript	128
TeX	51
R	42
CSS	39



tensorflow

Pull requests Issues Marketplace Explore



Repositories	31K
Code	4M
Commits	412K
Issues	56K
Topics	175
Wikis	3K
Users	420

Languages	
Python	17,069
Jupyter Notebook	6,421
HTML	653
C++	567
Java	409
Shell	304
JavaScript	233
TeX	68
R	64
C	62



Tensorflow

TensorFlow is an open source software library for numerical computation.

[See topic](#)

31,723 repository results

[Sort: Most stars ▾](#)

tensorflow/tensorflow



★ 95.7k

Computation using data flow graphs for scalable machine learning

[tensorflow](#) [python](#) [machine-learning](#)

Apache-2.0 license Updated 43 minutes ago

tensorflow/models



★ 32.7k

Models and examples built with *TensorFlow*

Apache-2.0 license Updated 3 hours ago

keras-team/keras



★ 28.1k

Deep Learning for humans

[deep-learning](#) [tensorflow](#) [python](#)

Updated 7 hours ago 10 issues need help

[Advanced search](#) [Cheat sheet](#)

Open Source Models

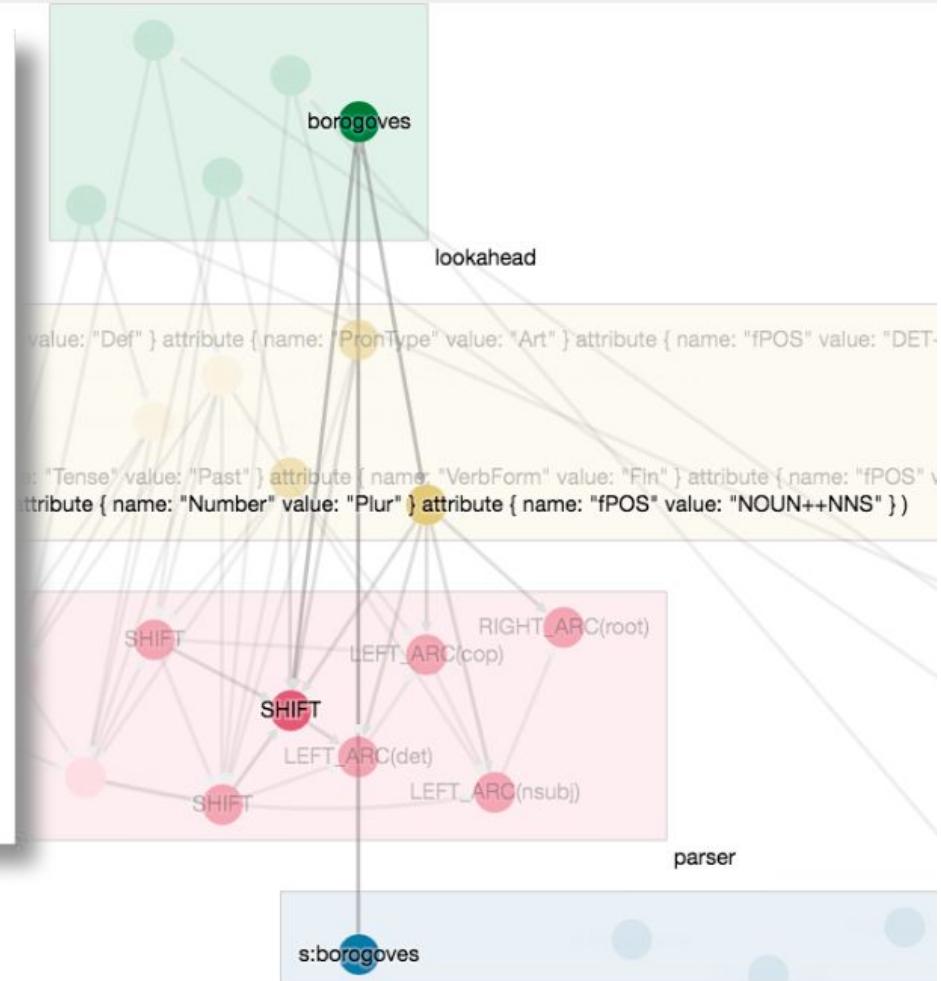
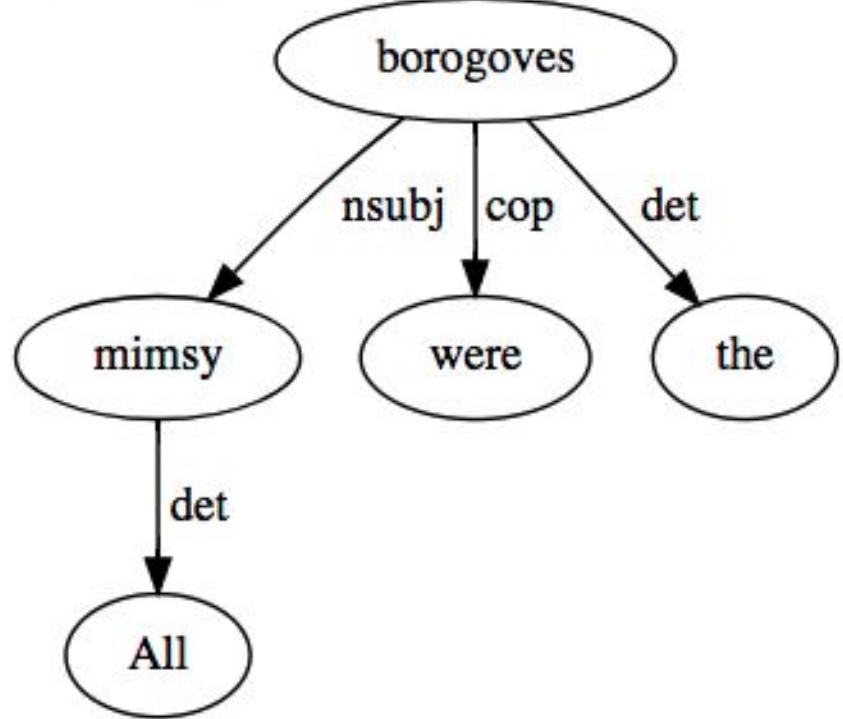
github.com/tensorflow/models

(many include trained weights, and
support transfer learning)

SyntaxNet and ParseySaurus

[research.googleblog.com/2017/03/an-upgrade-to-syntaxne
t-new-models-and.html](https://research.googleblog.com/2017/03/an-upgrade-to-syntaxnet-new-models-and.html)

Text: All mimsy were the borogoves



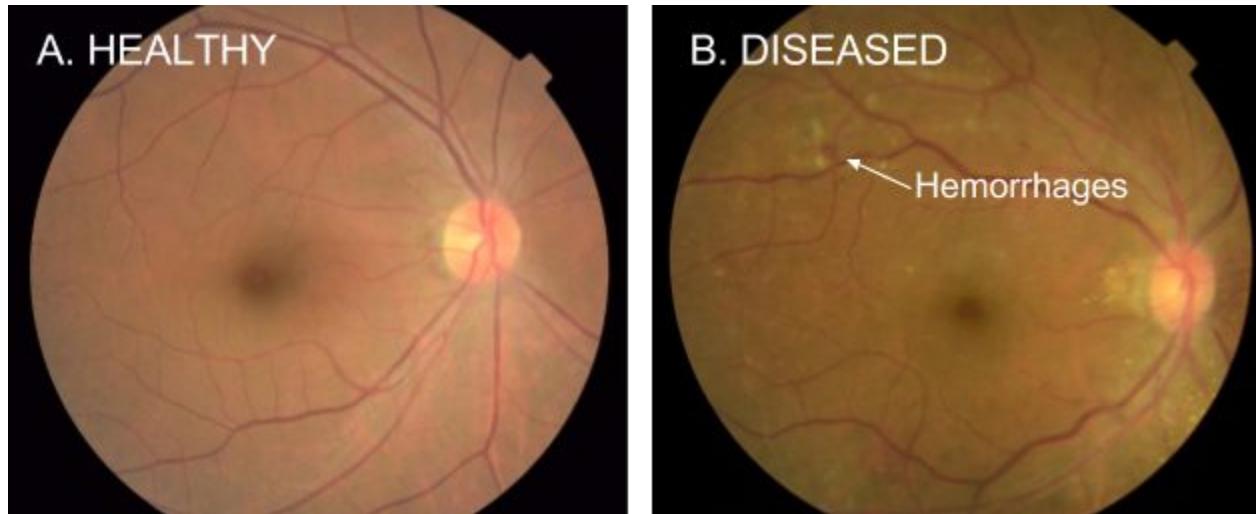
Inception: Image classification



An Alaskan Malamute (left) and a Siberian Husky (right). Images from Wikipedia.

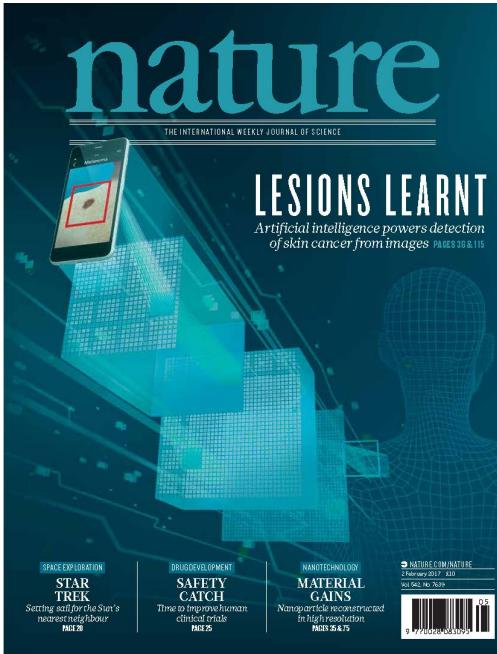
<https://research.googleblog.com/2016/08/improving-inception-and-image.html>

Detection of Diabetic Eye Disease



<https://research.googleblog.com/2016/11/deep-learning-for-detection-of-diabetic.html>

Skin Cancer Image Classification

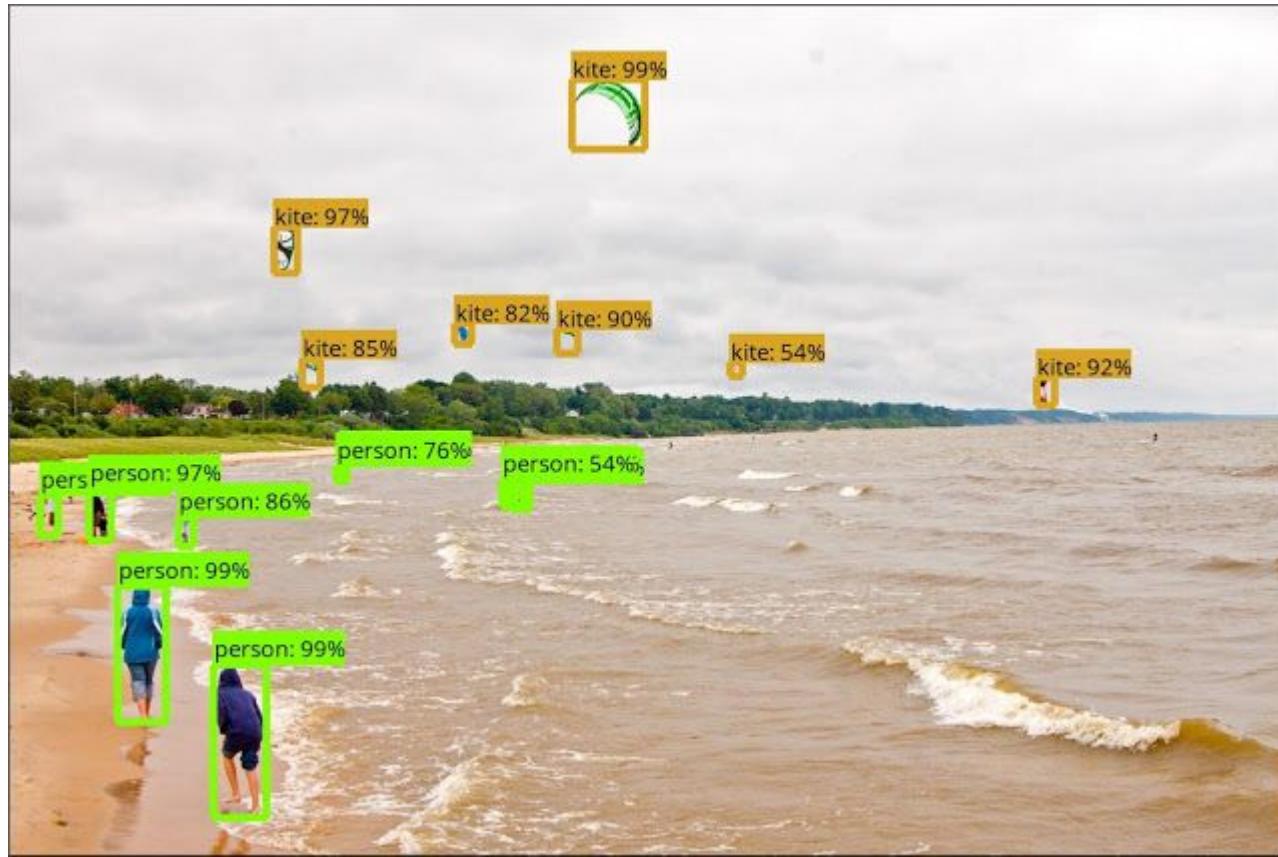


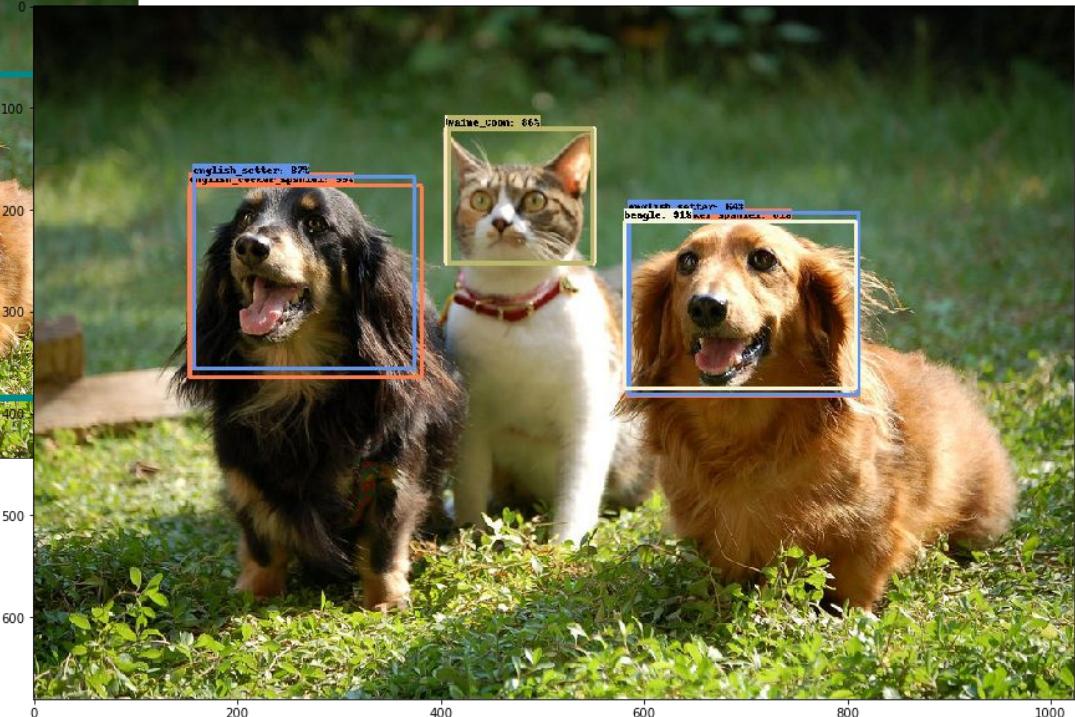
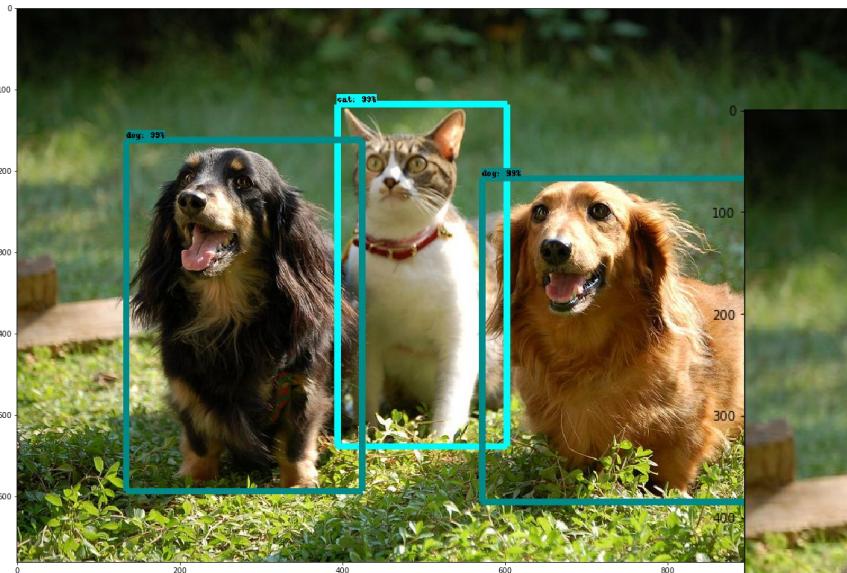
<http://www.nature.com/nature/journal/v542/n7639/full/nature21056.html>

TensorFlow Object Detection API

research.googleblog.com/2017/06/supercharge-your-computer-vision-models.html

TensorFlow Object Detection API





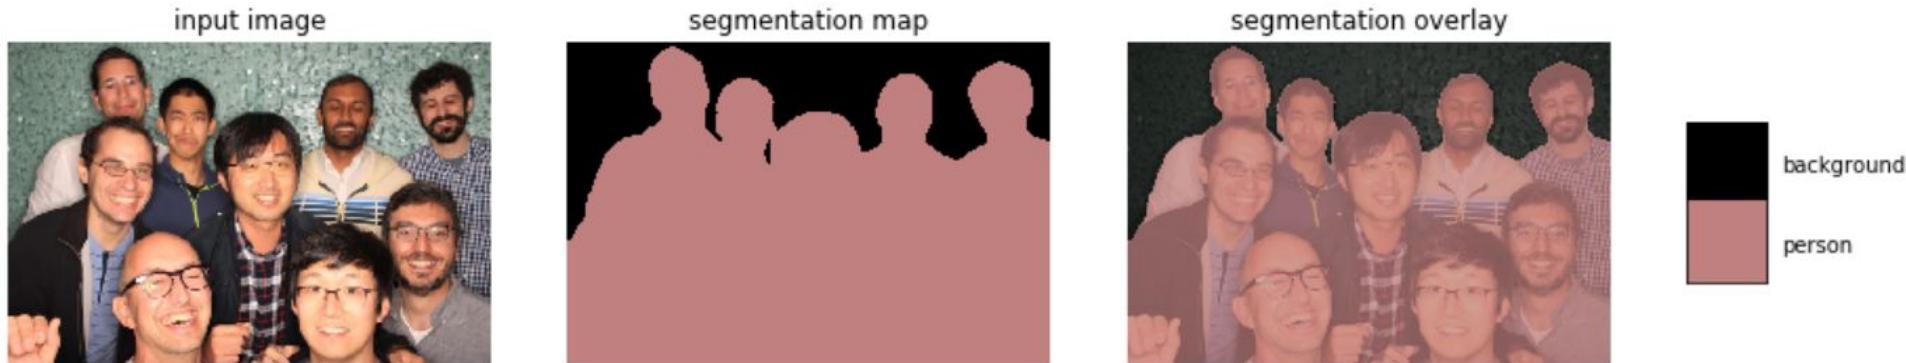
Object Detection API +Transfer Learning

bit.ly/tf-workshop-aicon

<https://goo.gl/images/xf45oH>, by
<https://www.flickr.com/photos/raneko/>

TensorFlow Object Detection API

Now with **more detection!**



<https://research.googleblog.com/2018/03/semantic-image-segmentation-with.html>

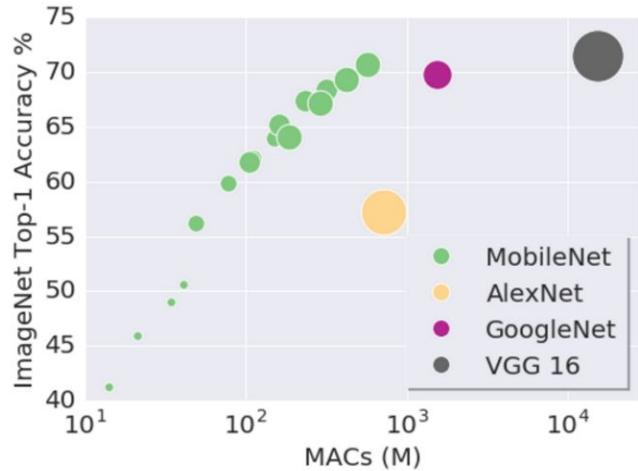
Real-time Video Segmentation (on mobile!)

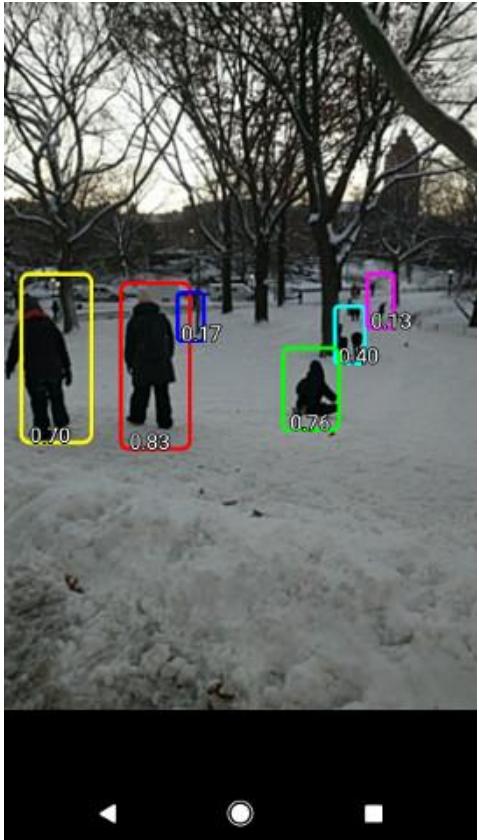
<https://research.googleblog.com/2018/03/mobile-real-time-video-segmentation.html>



MobileNets: Open-Source Models for Efficient On-Device Vision

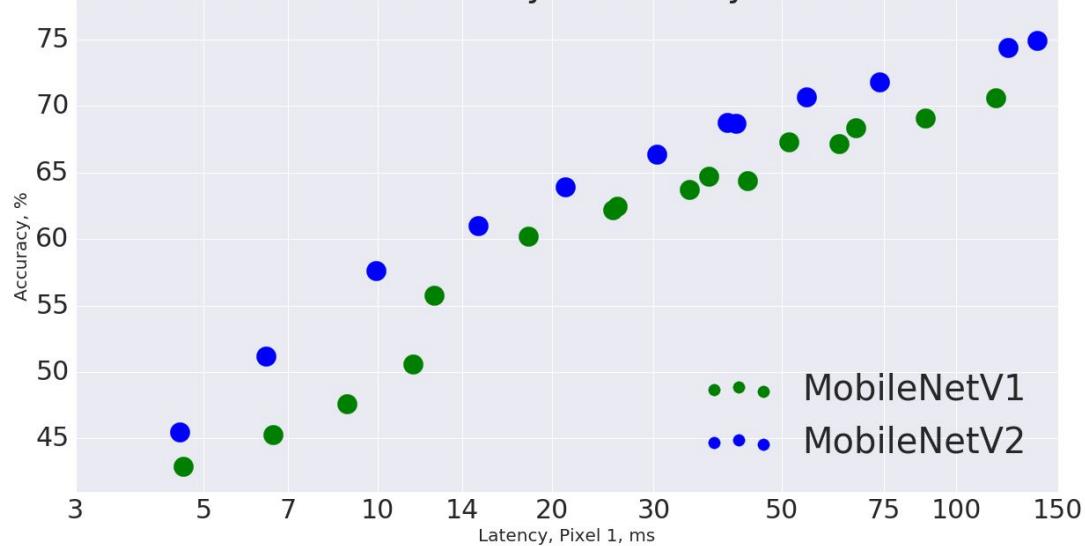
<https://research.googleblog.com/2017/06/mobilenets-open-source-models-for.html>





One of the newer Android demos,
[TF Detect](#), uses a **MobileNet**
model trained using the
Tensorflow Object Detection API.

Accuracy vs Latency



- **2x fewer operations**
- **30% fewer parameters**
- **30-40% faster** on a Google Pixel phone
- What about accuracy?
Higher!

MobileNet v2!

<https://research.googleblog.com/2018/04/mobilenetv2-next-generation-of-on.html>

Shared Research in TensorFlow

Inception: tensorflow.org/tutorials/image_recognition

Object Detection API: research.googleblog.com/2017/06/supercharge-your-computer-vision-models.html,
<https://cloud.google.com/blog/big-data/2017/09/performing-prediction-with-tensorflow-object-detection-models-on-google-cloud-machine-learning-engine>

ParseySaurus: research.googleblog.com/2017/03/an-upgrade-to-syntaxnet-new-models-and.html

MobileNets: <https://research.googleblog.com/2017/06/mobilenets-open-source-models-for.html>

Show and Tell <https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

Translation <https://research.googleblog.com/2016/09/a-neural-network-for-machine.html>

Summarization <https://research.googleblog.com/2016/08/text-summarization-with-tensorflow.html>

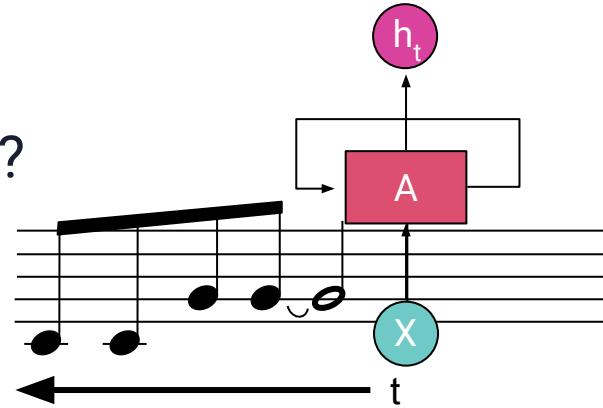
Pathology <https://research.googleblog.com/2017/03/assisting-pathologists-in-detecting.html>

Many more: <https://github.com/tensorflow/models/>

Project Magenta

Can ML generate compelling media?

- Music
- Images and Video
- Text (Jokes, Stories)
- Interestingness/Surprise
- Structure learning
- Measure success



magenta.tensorflow.org



Style Transfer



goo.gl/WPJtVe

TensorFlow and the OSS community

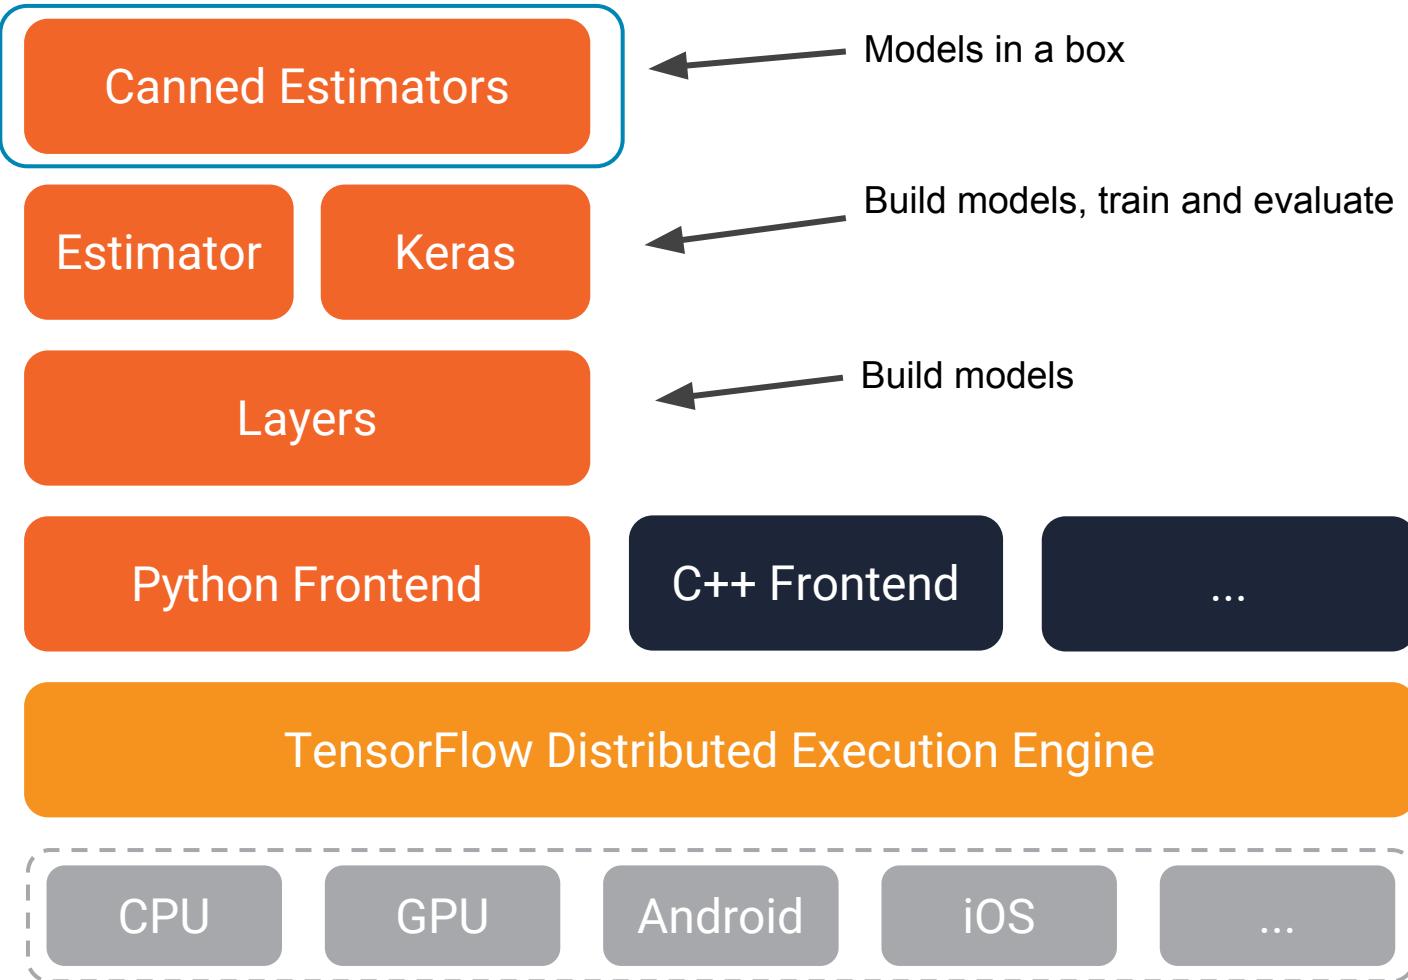


TensorFlow was open-sourced in large part to allow the community to improve it with contributions.

- The TensorFlow team has set up [processes](#) to manage pull requests, review and route issues filed, and answer [Stack Overflow](#) and [mailing list](#) questions.
- So far, we've had more than 790 external contributors add to the code, from small documentation fixes to large additions

TensorFlow's High-level APIs





fchollet/keras

Reference Implementation
TensorFlow backend
Multiple non-TF backends

tf.keras

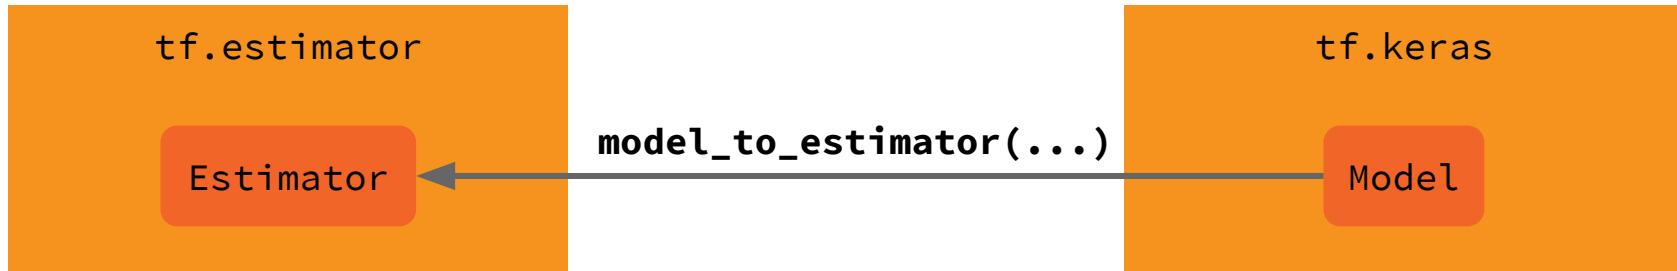
Custom TensorFlow backend
Integration with Estimators:
Distributed Execution
Integration with serving

Keras 2.0 API Spec

tf.layers and tf.keras

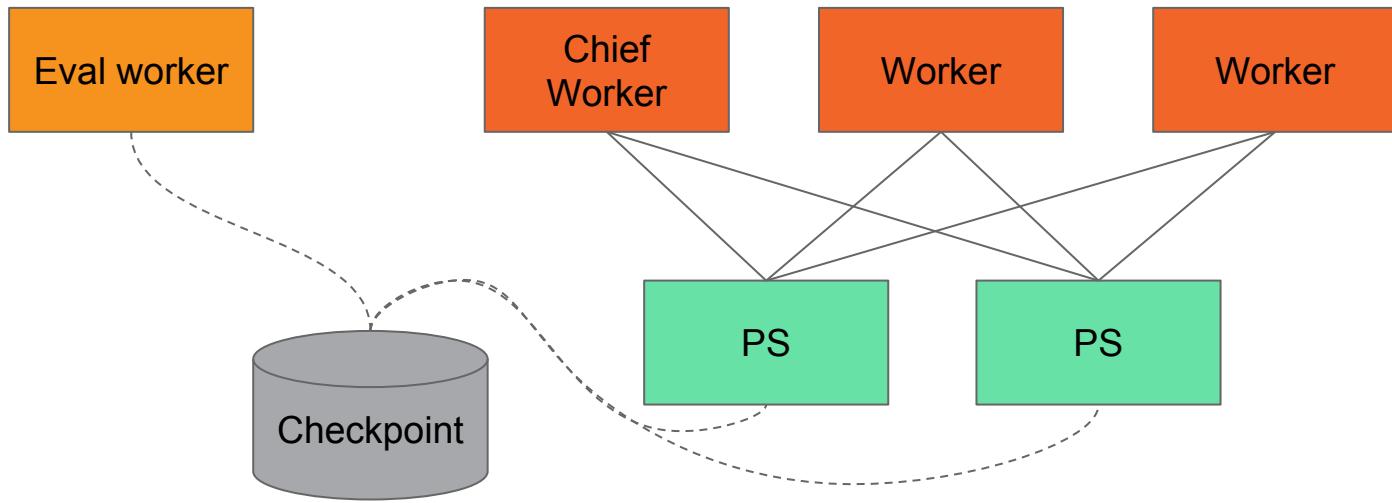
- tf.keras implements Keras API spec *exactly*
- tf.layers.Dense contains the implementation for tf.keras.Dense
- tf.keras is as thin a wrapper as possible

Keras/Estimator Integration



Estimators

- Implement:
 - training loop
 - checkpointing
 - Built-in integration with TensorBoard
 - Export to tensorflow/serving
- Supports complex use cases, custom inference, evaluation, etc.
- Complete implementations of standard models (linear, DNN, ...)



Estimators

Ready for distributed execution

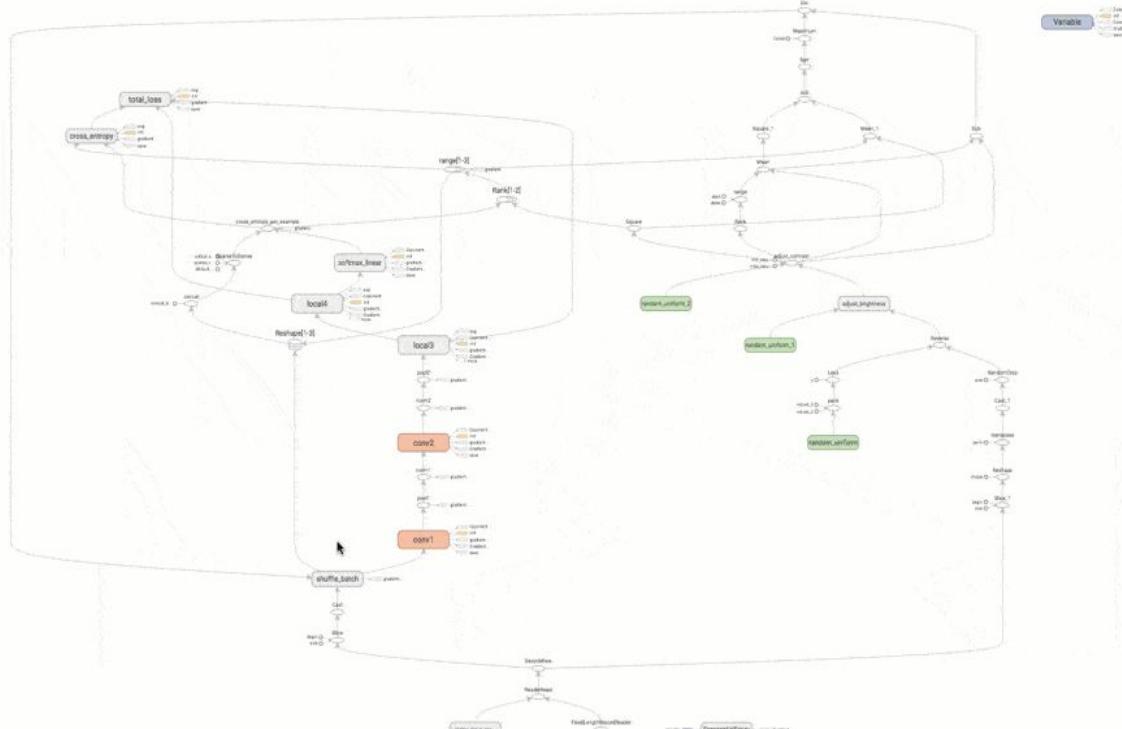
- Recovery after worker failure
- Distributed device placement
- Accelerator support

Run local or distributed; run on different devices, without changing your model.

The TensorFlow ecosystem...



Main Graph

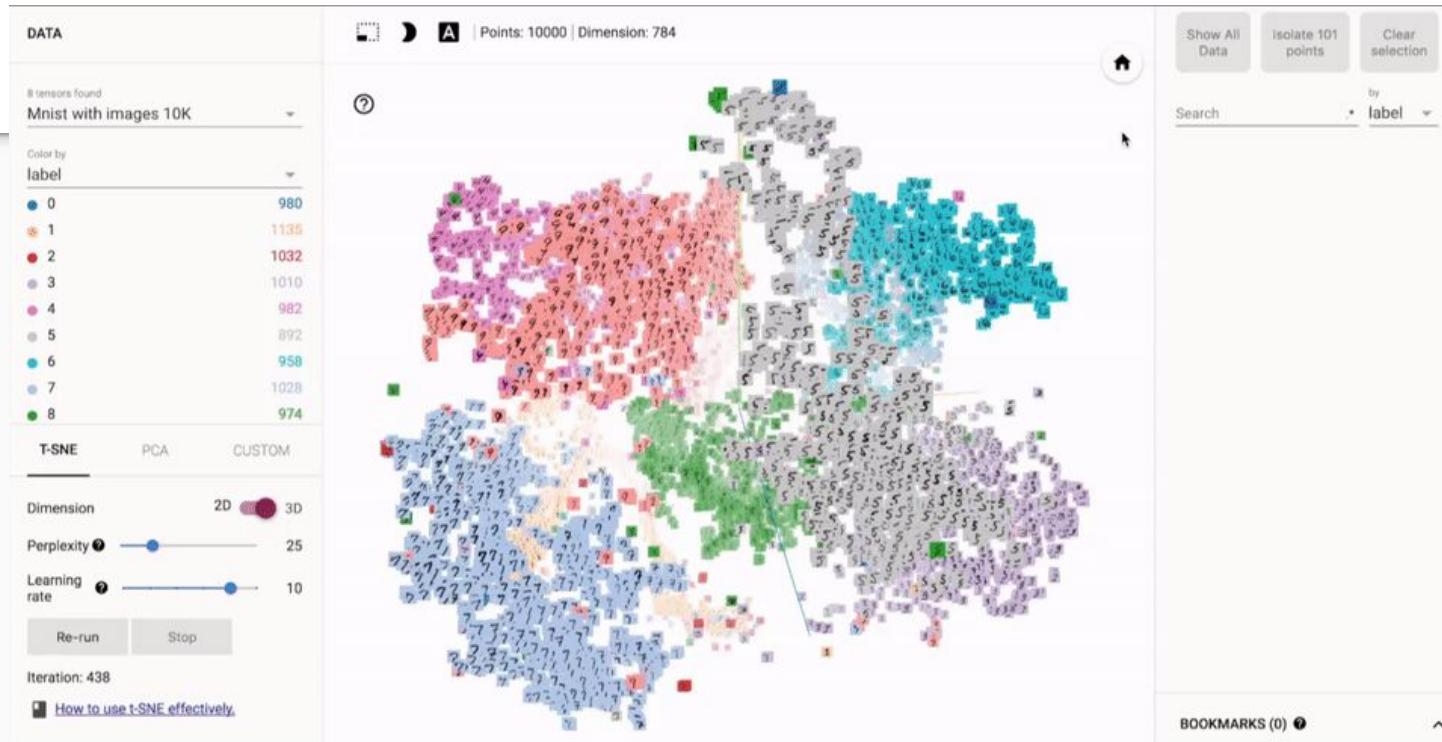


Auxiliary nodes



- Dataflow edge
- ↔ Control dependency edge
- ↔ Reference edge

TensorBoard



https://www.tensorflow.org/versions/r0.12/how_tos/embedding_viz/index.html



Iterations

000,815

Learning rate

0.03

Activation

Sigmoid

Regularization

None

Regularization rate

0

Problem type

Classification

DATA

Which dataset do you want to use?



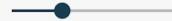

Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

X_1

X_2

X_1^2

X_2^2

$X_1 X_2$

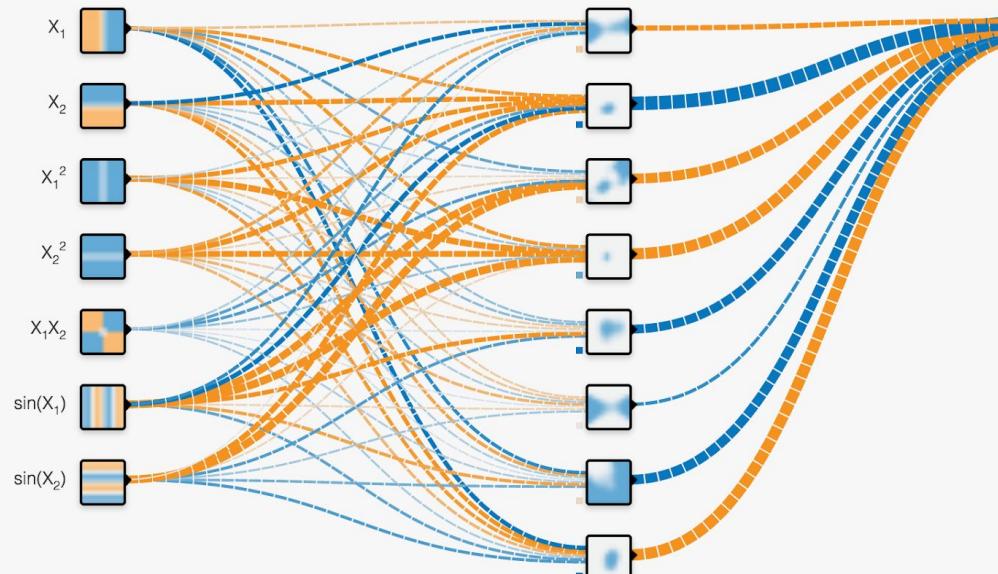
$\sin(X_1)$

$\sin(X_2)$

1 HIDDEN LAYER

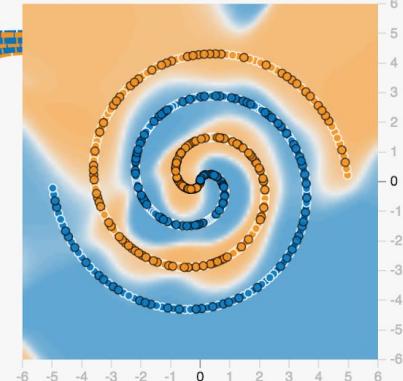
+ -

8 neurons



OUTPUT

Test loss 0.037
Training loss 0.026



Colors shows data, neuron and weight values.



Show test data

Discretize output

Using Estimators: LinearClassifier and DNNClassifier



These slides: bit.ly/tf-workshop-aicon

Create a Kaggle.com account

Go to kaggle.com/yufengg/kernels

OR

Install TensorFlow & Jupyter: www.tensorflow.org/install/

pip install tensorflow pandas jupyter

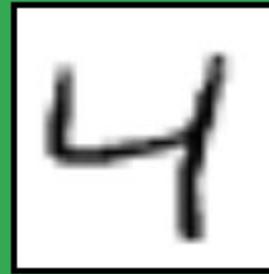
Clone GitHub Repo: bit.ly/tensorflow-workshop

Optional: download the ‘Fashion-MNIST’ files:
github.com/zalandoresearch/fashion-mnist#get-the-data

Estimator workflow

```
model =  
    tf.estimator.LinearClassifier(feature_columns, ... )  
  
model.train(train_input_fn, steps)  
  
model.evaluate(eval_input_fn)  
  
model.predict(predict_input_fn)  
  
model.export_savedmodel(serving_input_receiver_fn)
```

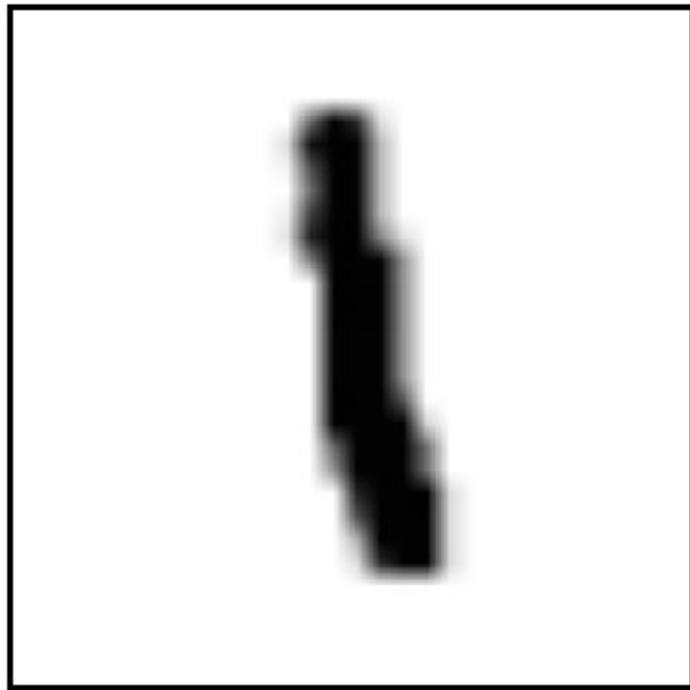
MNIST: The “Hello World” Of ML



fashion-MNIST: Making MNIST fashionable again

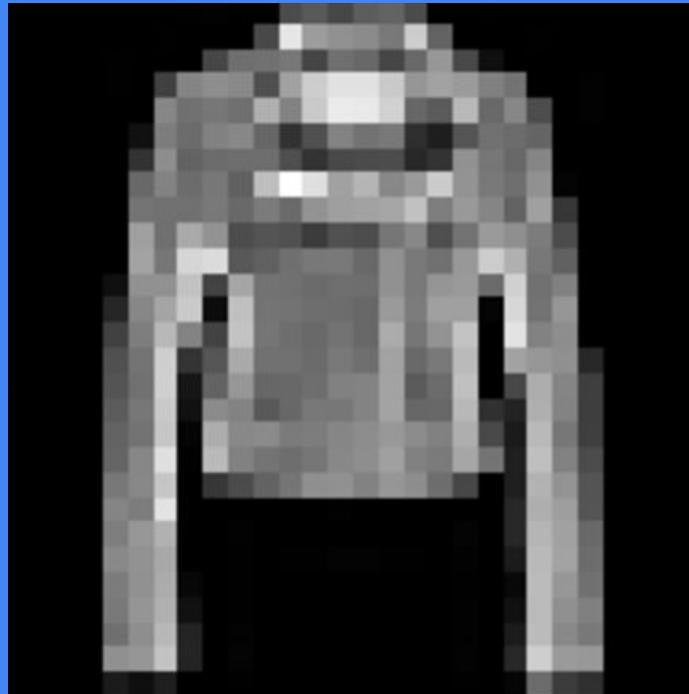


Computer Vision -- MNIST



2

Computer Fashion -- fashion-MNIST



Lab: using TensorFlow's Estimator APIs

Fork my notebook:
kaggle.com/yufengg/fashion-mnist

Wait, how do I fork a notebook???



Fit to screen



Download PNG



Run model_14841...



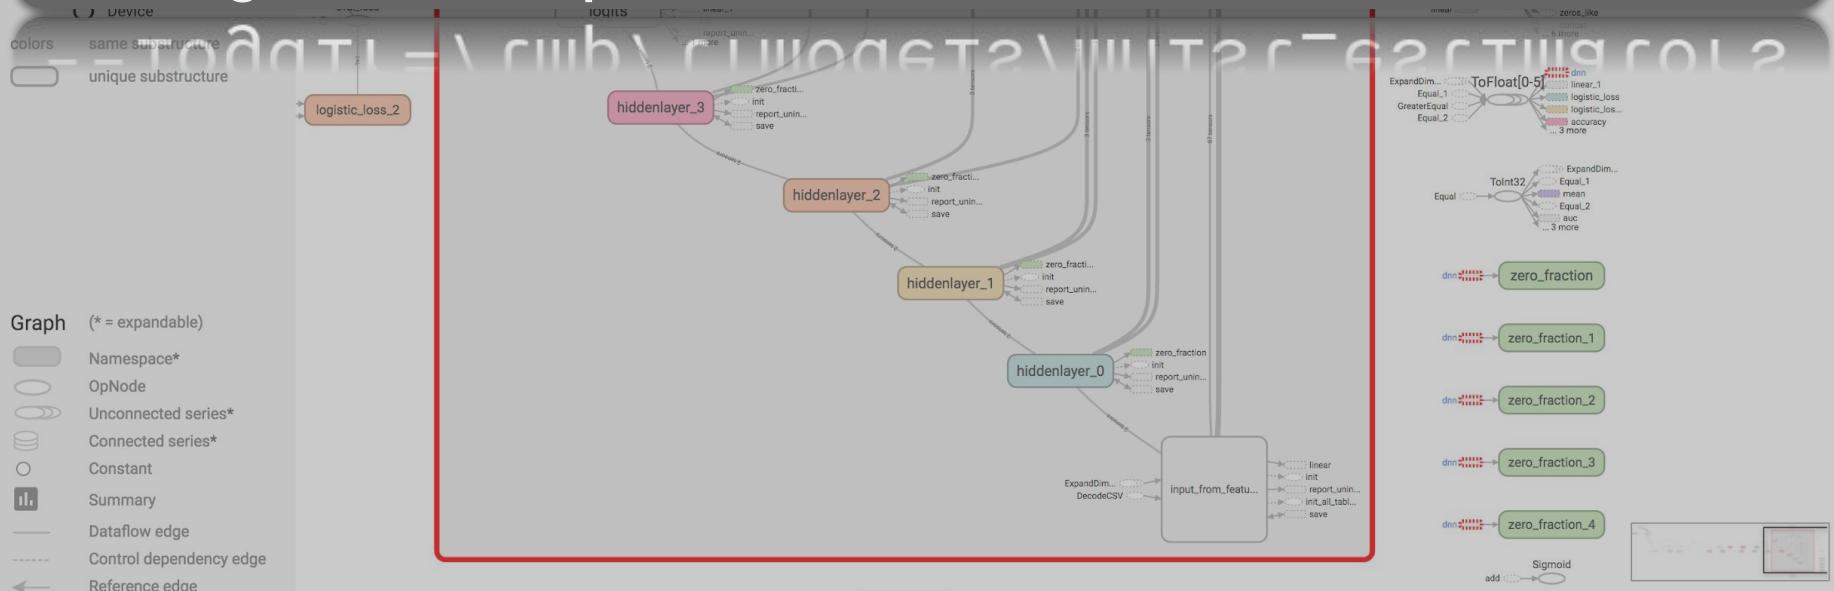
(18)

Session

Run

18

tensorboard --logdir=/tmp/tfmodels/mnist_estimators

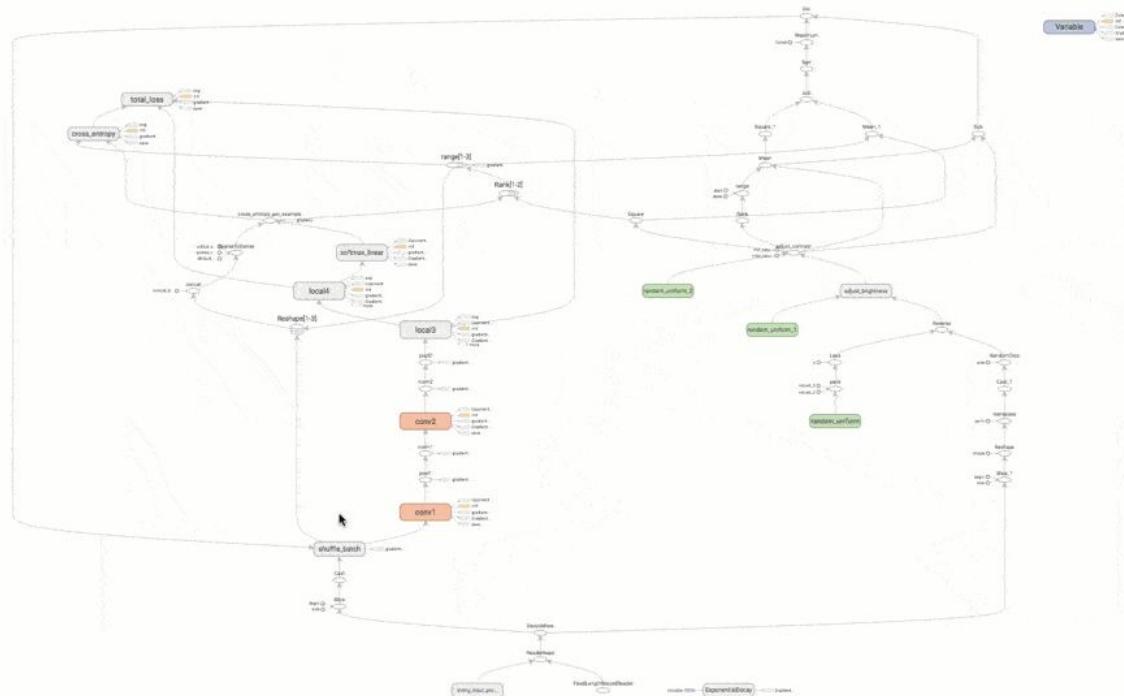


Fit to screen
Run cifar-train
Choose File

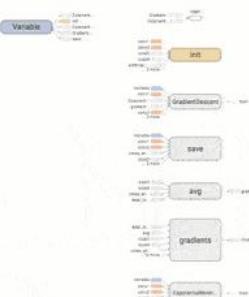
Upload Choose File

Color Structure
color: same substructure
gray: unique substructure

Main Graph



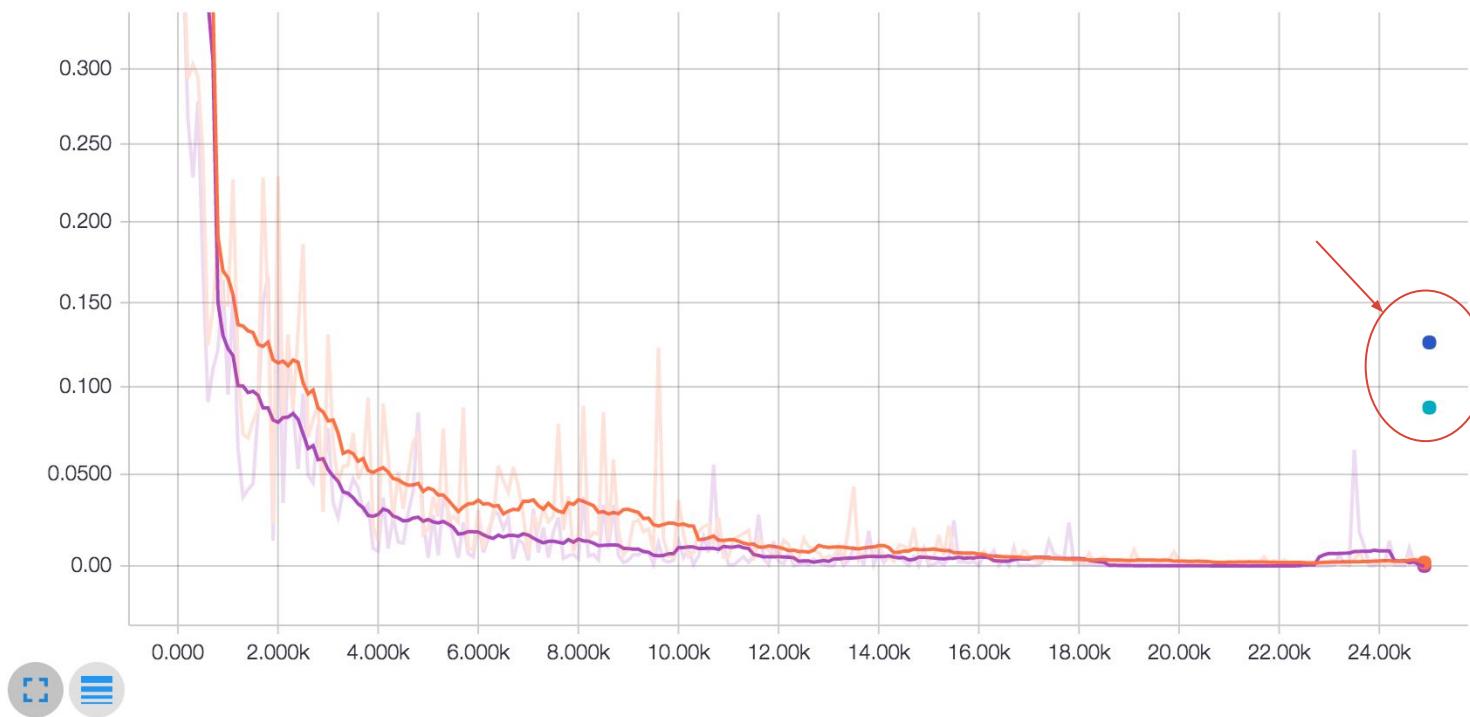
Auxiliary nodes



Graph (* = expandable)

- Namespace*
- OpNode
- Unconnected series*
- Connected series*
- Constant
- Summary
- Dataflow edge
- Control dependency edge
- Reference edge

loss



Comparing two different optimizers for DNNClassifier

- Show data download links
 Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing



Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs

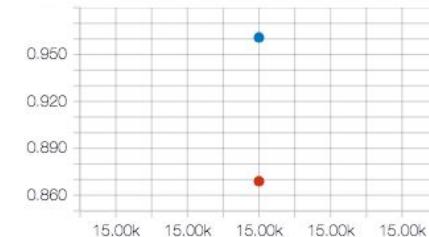
- fashion/.
 fashion/eval
 reg/.
 reg/eval

Q.*

Tags matching ./*/(all tags)

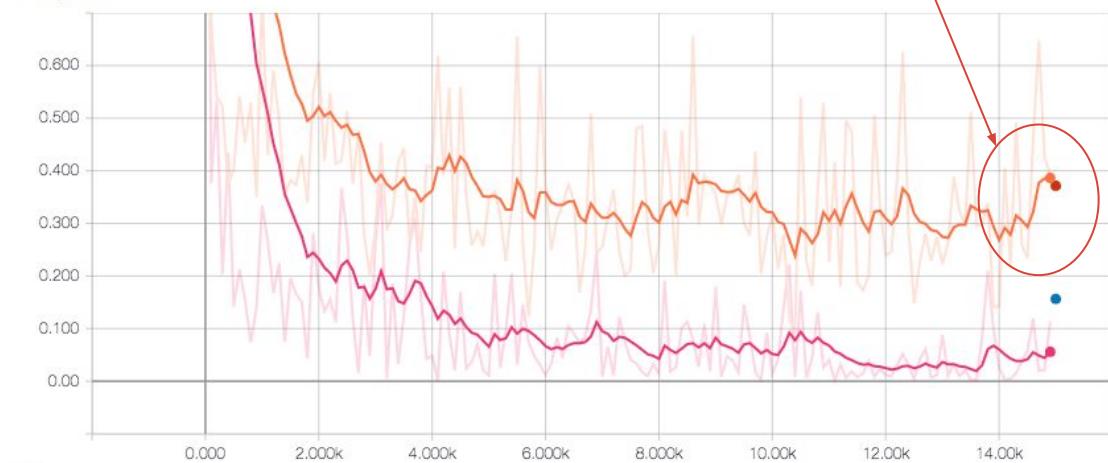
8

accuracy



fashion-mnist

average_loss

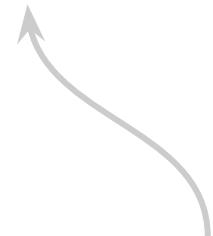


What's TensorFlow made of?

A multidimensional array.



TensorFlow

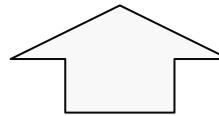


A graph of operations.

Tensors - generalized matrices

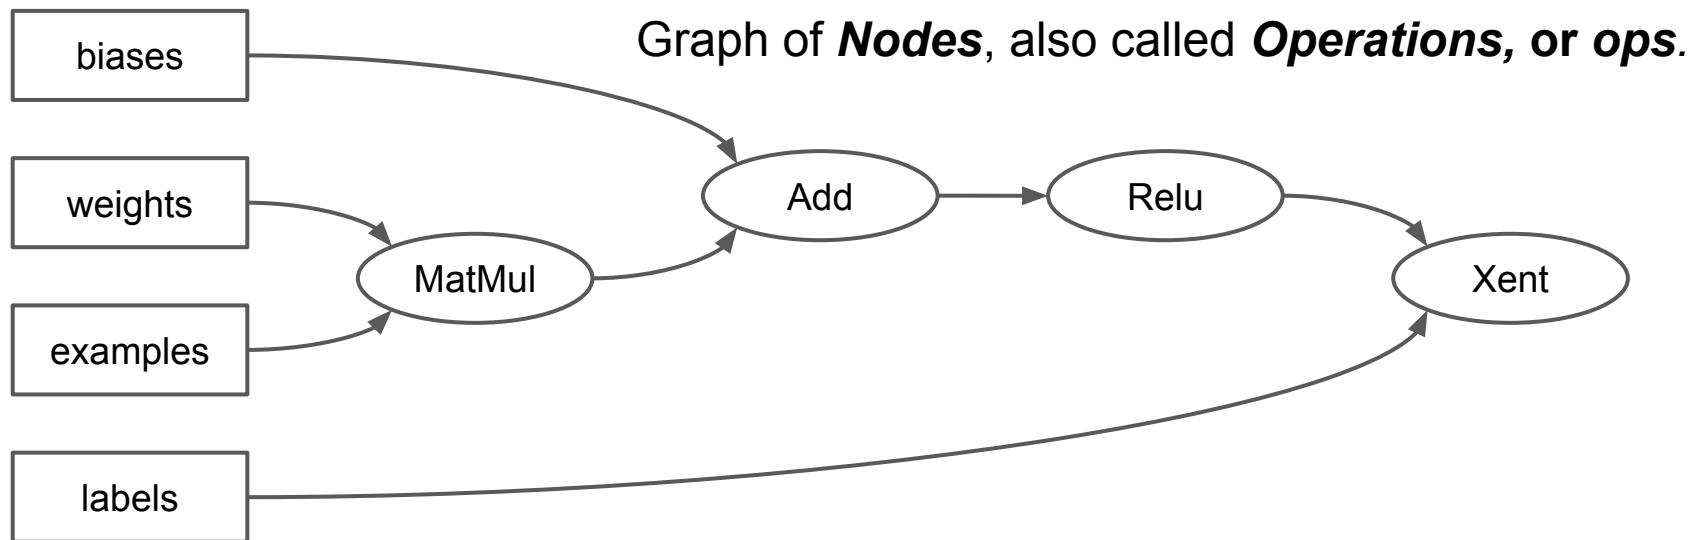
Tensors have a Shape that's described with a *vector*.

[10000, 256, 256, 3]



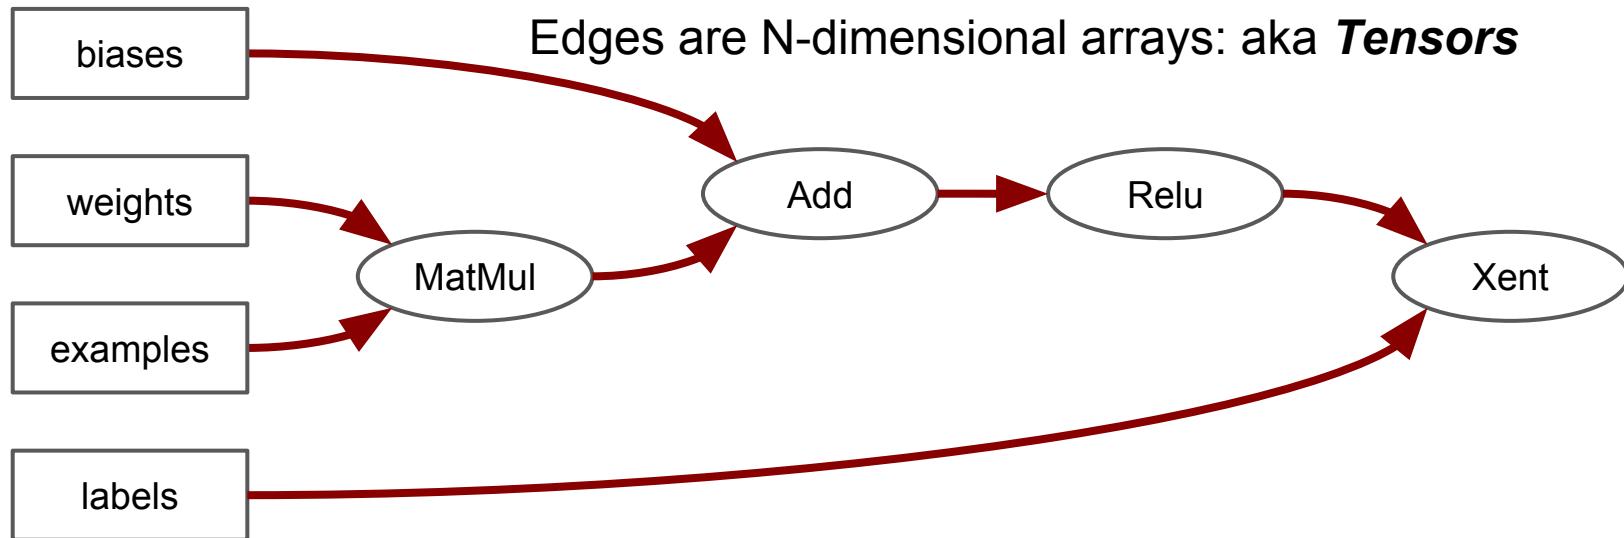
- **10000** Images
- Each Image has **256 Rows**
- Each Row has **256 Pixels**
- Each Pixel has **3 channels** (RGB)

Computation is a dataflow graph



Computation is a dataflow graph

with tensors



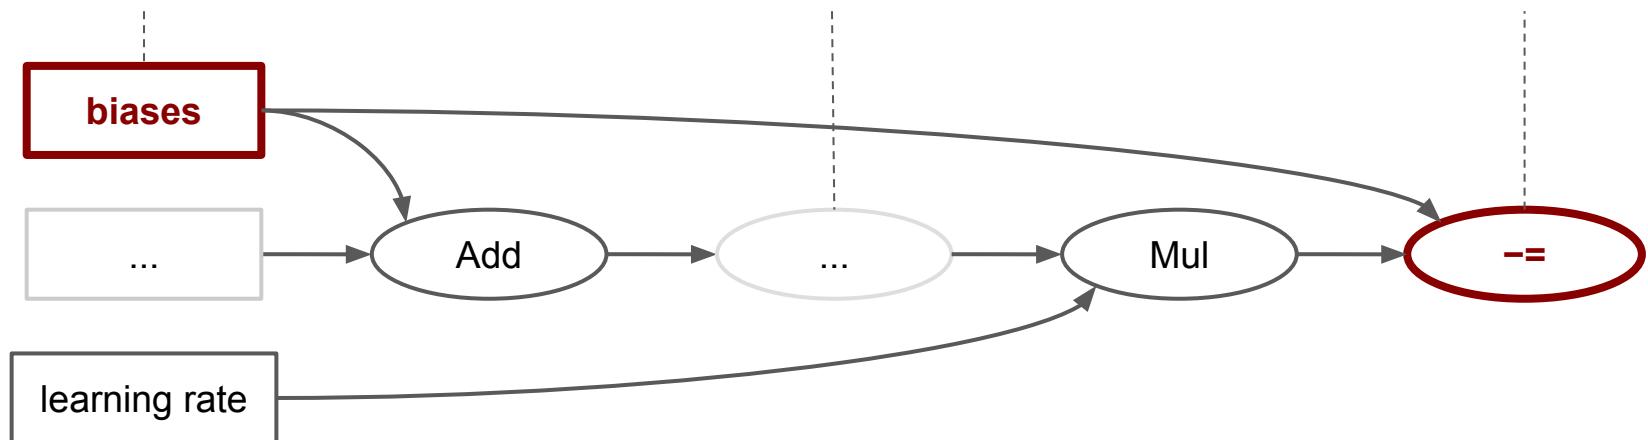
Computation is a dataflow graph

with state

'Biases' is a **variable**

Some ops compute **gradients**

$\text{--} =$ updates **biases**





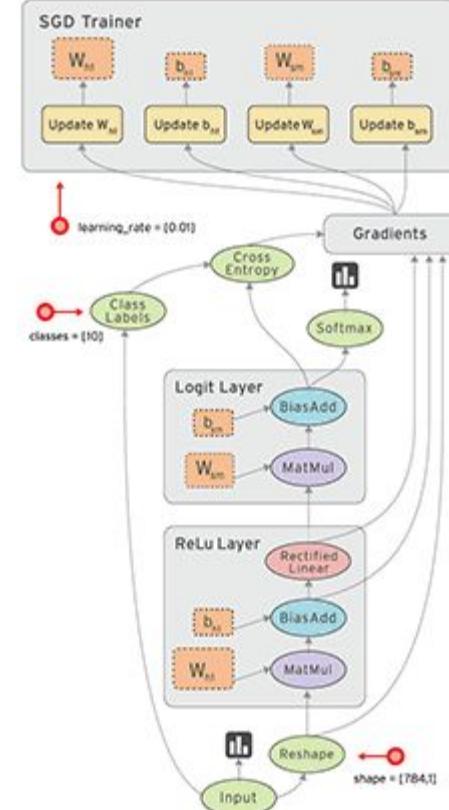
A multidimensional array.

A graph of operations.



Operates over **tensors**: *n-dimensional arrays*
Using a **flow graph**: *data flow computation framework*

- Flexible, intuitive construction
- Support for threads, queues, and asynchronous computation; **distributed runtime**
- **automatic differentiation**
- Train on CPUs, GPUs, ...and coming soon, **TPUS...**
- Run wherever you like





Building a computation graph, then executing it:

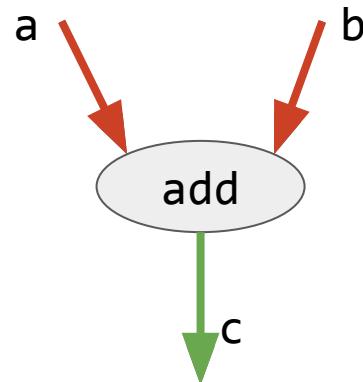
- **Portability:** the graph can be executed immediately, or saved to use later, and it can run on multiple platforms: CPUs, GPUs, TPUs, mobile, embedded.
 - Deploy to production without having to depend on any of the code that built the graph, only the runtime necessary to execute it.
- **Transformable and Optimizable:**
 - produce a more optimal version of the graph for a given platform.
 - perform memory or compute optimizations, and trade off between them.
- Support for **distributed execution**.

Build a graph; then run it.

```
...  
# build graph  
c = tf.add(a, b)
```

```
...
```

```
with tf.Session() as session:  
    # run graph  
    value_of_c = session.run(c, {a=1, b=2})
```



tf.nn

[Overview](#)[all_candidate_sampler](#)
[atrous_conv2d](#)
[atrous_conv2d_transpose](#)
[avg_pool](#)
[avg_pool3d](#)
[batch_normalization](#)
[batch_norm_with_global_normalization](#)[bias_add](#)
[bidirectional_dynamic_rnn](#)[compute_accidental_hits](#)[conv1d](#)[conv2d](#)[conv2d_backprop_filter](#)[conv2d_backprop_input](#)[conv2d_transpose](#)[conv3d](#)[conv3d_backprop_filter_v2](#)[conv3d_transpose](#)[convolution](#)[crelu](#)[ctc_beam_search_decoder](#)[ctc_greedy_decoder](#)[softmax](#)Module `tf.nn`Defined in [tensorflow/python.ops/nn.py](#)

Neural network support.

[See the TensorFlow guide](#)

Even though TF has Neural Network ops... directly using these basic building blocks can be pretty low-level

[atrous_conv2d\(...\)](#) : Atrous convolution (a.k.a. convolution with holes or dilated convolution).

[atrous_conv2d_transpose\(...\)](#) : The transpose of `atrous_conv2d`.

[avg_pool\(...\)](#) : Performs the average pooling on the input.

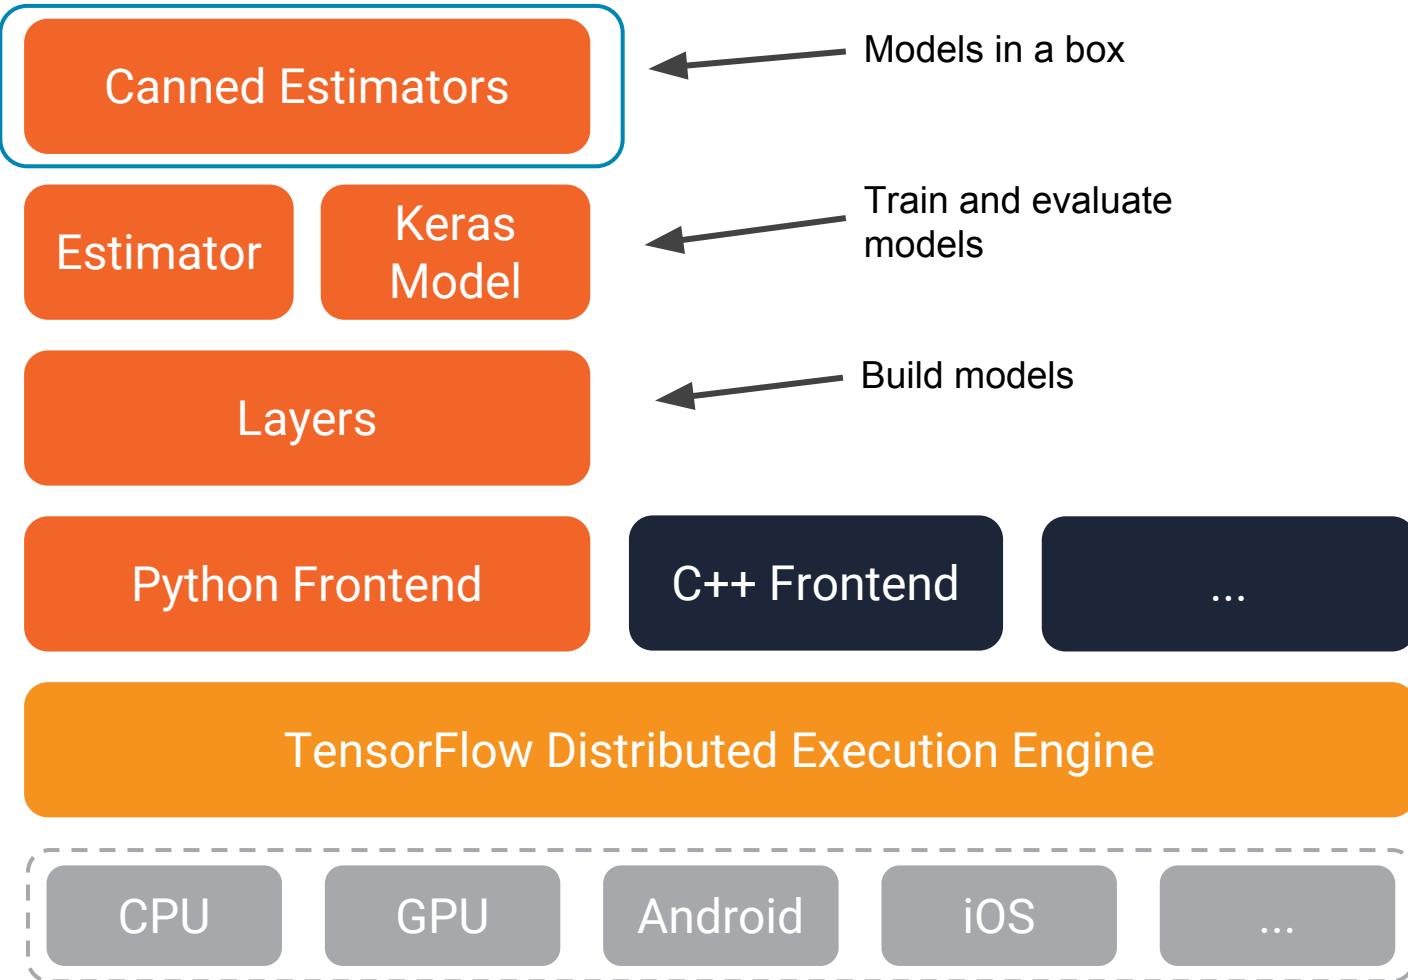
[avg_pool3d\(...\)](#) : Performs 3D average pooling on the input.

[batch_norm_with_global_normalization\(...\)](#) : Batch normalization.

[batch_normalization\(...\)](#) : Batch normalization.

[bias_add\(...\)](#) : Adds bias to value.

[bidirectional_dynamic_rnn\(...\)](#) : Creates a dynamic version of bidirectional recurrent neural network.

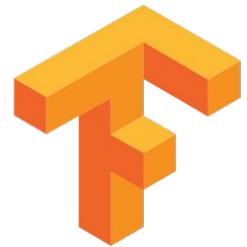


Many levels of abstraction



Choose the right one for you:

- Layers, losses, metrics
- Training/Eval loop functions
- Estimator (BaseEstimator)
 - Any model you want, but must separate input from the rest of the model
- Predefined estimators
 - `LinearClassifier`, `DNNClassifier`,
`DNNRegressor`, ... **DNNLinearCombinedClassifier**
 - Limited configuration options: feature columns, metrics.



These slides: bit.ly/tf-workshop-aicon

Create a Kaggle.com account

Go to kaggle.com/yufengg/kernels

OR

Install TensorFlow & Jupyter: www.tensorflow.org/install/

pip install tensorflow pandas jupyter

Clone GitHub Repo: bit.ly/tensorflow-workshop

Optional: download the ‘Fashion-MNIST’ files:
github.com/zalandoresearch/fashion-mnist#get-the-data

Using the TensorFlow High-level APIs with tf.data

kaggle.com/yufengg/mushrooms-with-datasets-and-estimators

Too much to type?
Try this instead

kaggle.com/yufengg/kernels

But first, let's look at Facets

The Dataset interface

Data sources and functional transformations

Create a Dataset from one or more `tf.Tensor` objects:

```
Dataset.from_tensors((features, labels))
```

```
Dataset.from_tensor_slices((features, labels))
```

```
TextLineDataset(filenames)
```

The Dataset interface

Data sources and functional transformations

Or create a Dataset from another Dataset:

```
dataset.map(lambda x: tf.decode_jpeg(x))
```

```
dataset.repeat(NUM_EPOCHS)
```

```
dataset.batch(BATCH_SIZE)
```

...and many more.

The Dataset interface

Data sources and functional transformations

Or create a Dataset from a Python generator:

```
def generator():
    while True:
        yield ...
```

```
dataset.from_generator(generator, tf.int32)
```

Let's go to the land of the mushrooms!

kaggle.com/yufengg/mushrooms-with-datasets-and-estimators

“Wide & Deep”: Using the TensorFlow High-level APIs and feature engineering

Motivation - a "magical" food app



Launch and Iterate!



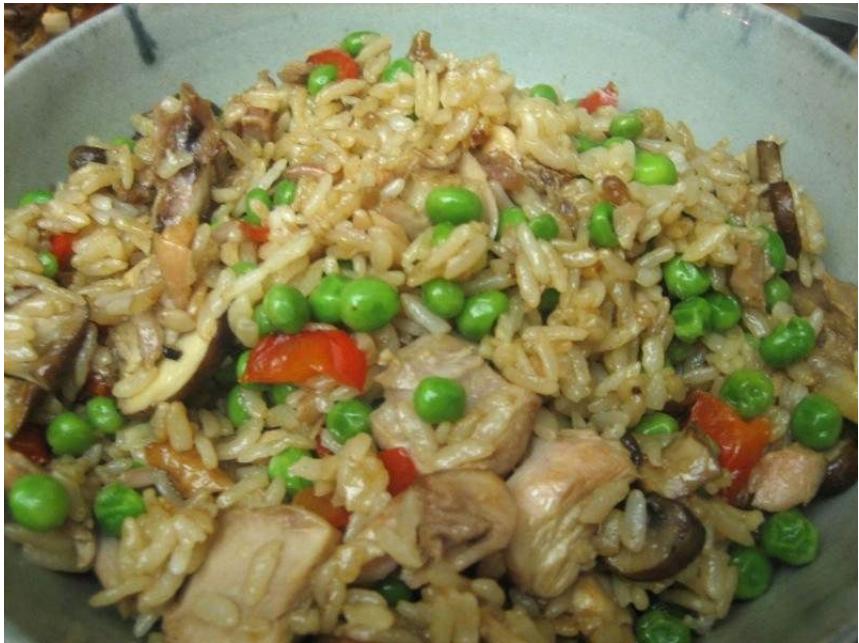
<https://upload.wikimedia.org/wikipedia/commons/2/2c/Fried-Chicken-Set.jpg>

Launch and Iterate!



https://upload.wikimedia.org/wikipedia/commons/f/fd/Chicken_and_Waffles_201_-_Evan_Swigart.jpg

Launch and Iterate!



... naive character matching

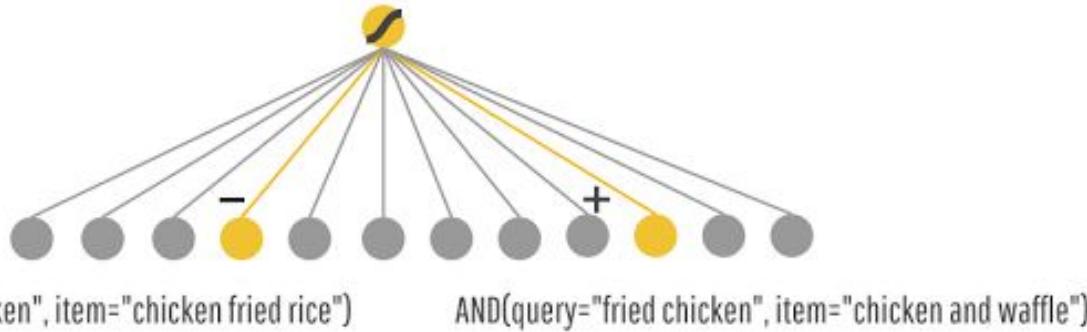


<https://www.flickr.com/photos/elsiehui/9575196704>

<https://www.flickr.com/photos/patrickwoodward/2208031777>

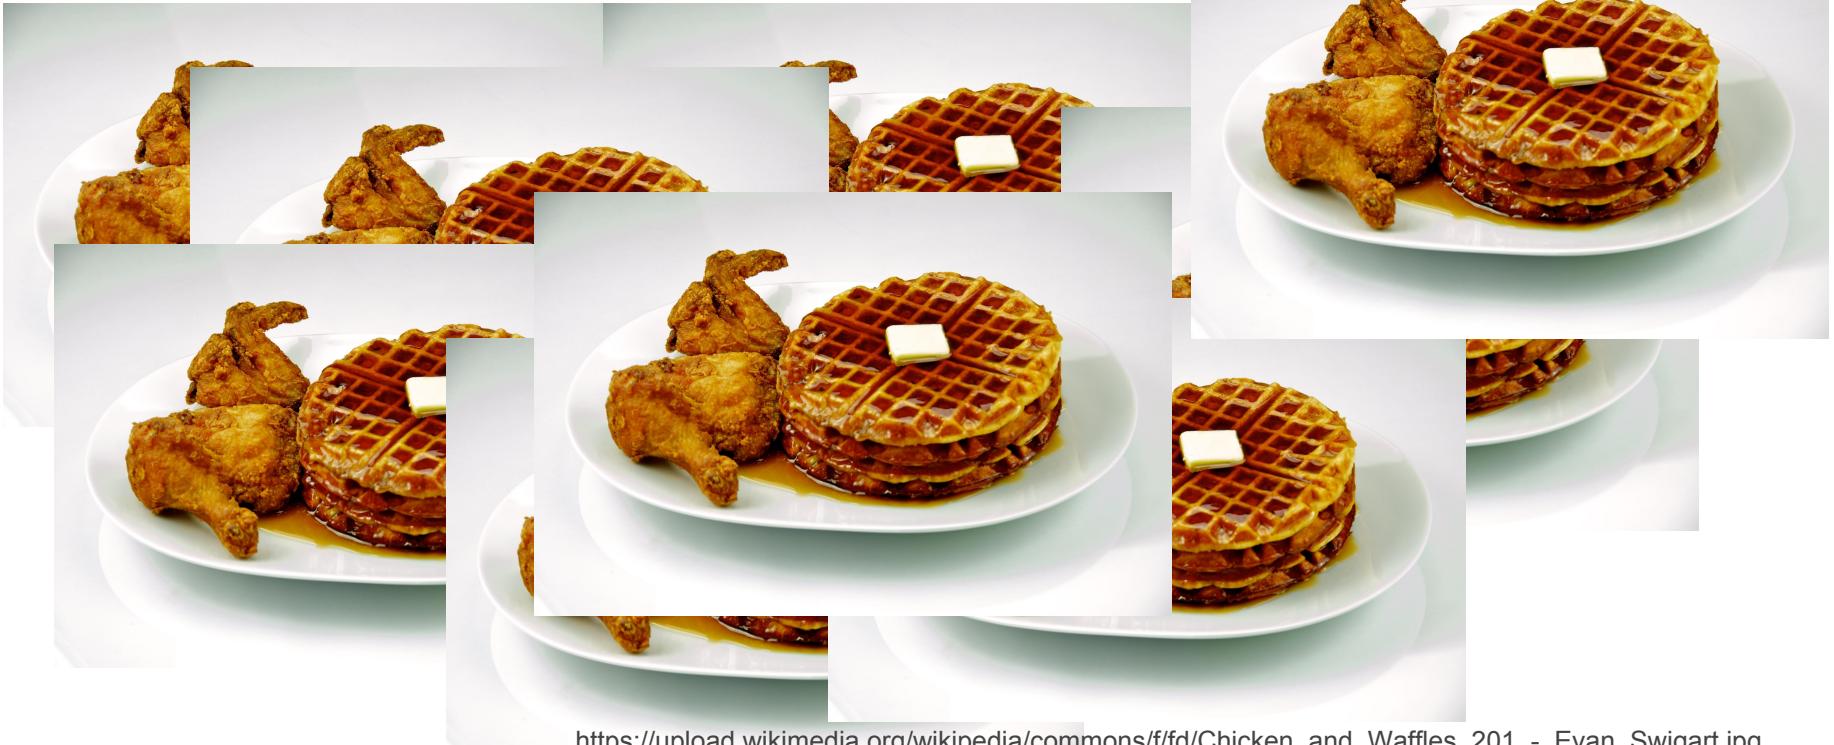
v2.0: memorize all the things

- Train a linear TF model



- Your app is gaining traction!

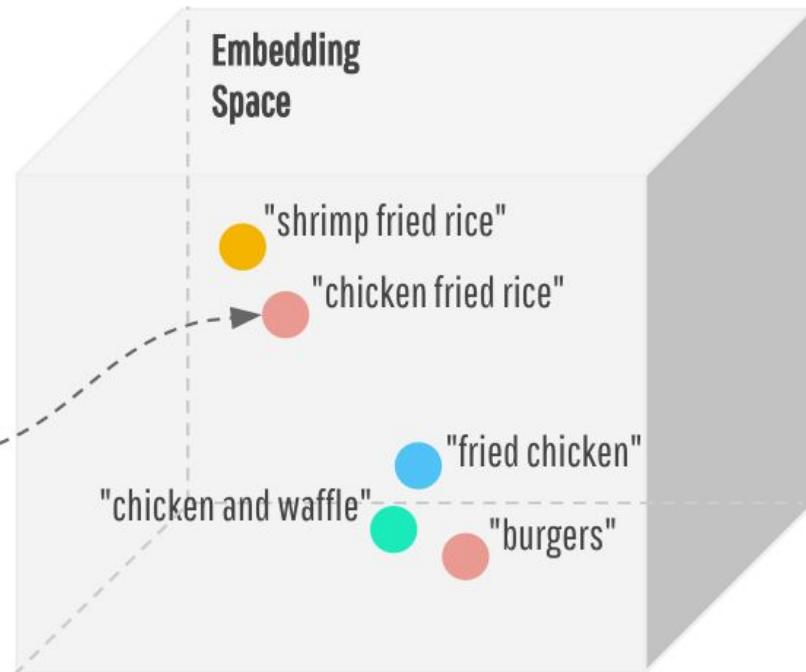
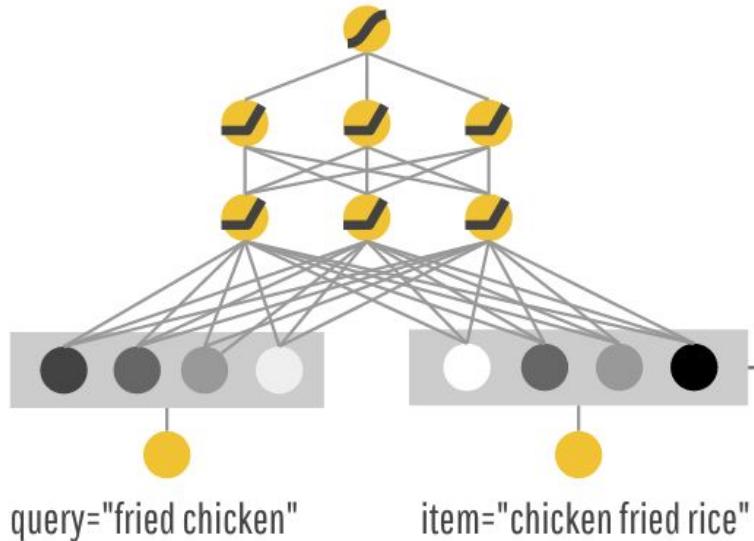
Problem: Your users are bored!



https://upload.wikimedia.org/wikipedia/commons/f/fd/Chicken_and_Waffles_201_-_Evan_Swigart.jpg

bit.ly/tf-workshop-aicon

v3.0: More generalized recommendations for all



No good deed goes unpunished

- Some recommendations are "too general"
 - Irrelevant dishes are being sent
- Your users are still picky 

No good deed goes unpunished



!=



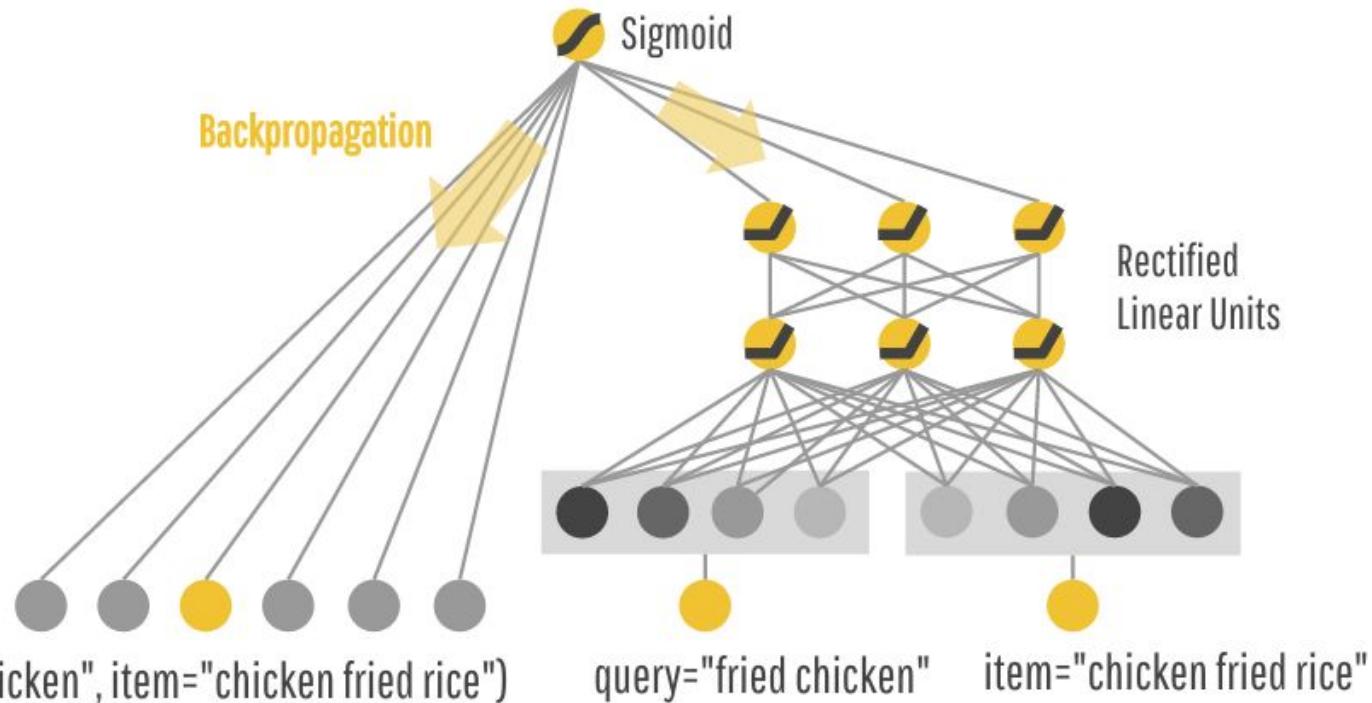
hot latte with whole milk

iced decaf latte with nonfat milk

Two types of requests:
specific and general

How to balance this?

v4.0: Why not both?



Wide & Deep

memorization
relevance

generalization
diversity

Lab: Wide and Deep: Using TensorFlow's high-level APIs

kaggle.com/yufengg/census-with-w-d-estimators



These slides: bit.ly/tf-workshop-aicon

Create a Kaggle.com account

Go to kaggle.com/yufengg/kernels

OR

Install TensorFlow & Jupyter: www.tensorflow.org/install/

pip install tensorflow pandas jupyter

Clone GitHub Repo: bit.ly/tensorflow-workshop

Optional: download the ‘Fashion-MNIST’ files:
github.com/zalandoresearch/fashion-mnist#get-the-data

Meet our dataset: US Census Data

- Just as exciting as chicken and waffles
- **Task:** predict the probability that the individual has an annual income of over 50K dollars
- Over 32k examples
- Extracted from the 1994 US Census by Barry Becker.

Meet our dataset: US Census Data

Column Name	Type	Description
age	Continuous	The age of the individual
workclass	Categorical	The type of employer the individual has (government, military, private, etc.).
fnlwgt	Continuous	The number of people the census takers believe that observation represents (sample weight). This variable will not be used.
education	Categorical	The highest level of education achieved for that individual.
education_num	Continuous	The highest level of education in numerical form.
marital_status	Categorical	Marital status of the individual.

Meet our dataset: US Census Data

Column Name	Type	Description
occupation	Categorical	The occupation of the individual.
relationship	Categorical	Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
race	Categorical	White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
sex	Categorical	Female, Male.
capital_gain	Continuous	Capital gains recorded.
capital_loss	Continuous	Capital Losses recorded.

Meet our dataset: US Census Data

Column Name	Type	Description
hours_per_week	Continuous	Hours worked per week.
native_country	Categorical	Country of origin of the individual.
income_bracket	Categorical	">50K" or "<=50K", meaning whether the person makes more than \$50,000 annually.

The return of Facets

Typical structure



- Load data
- Set up feature columns
- Create your model
- Run the training loop (fit the model)
- Evaluate your model's accuracy (and other metrics)
- (optional) Predict new examples



Estimators high-level structure

```
feature_columns = tf.feature_column.numeric_column(  
    "pixels", shape=784)  
model = tf.estimator.DNNClassifier(  
    [layer2_hidden_units, layer1_hidden_units],  
    feature_columns=feature_columns,  
    n_classes=2  
)  
model.train(input_fn=training_input_fn, steps=1000)  
model.evaluate(input_fn=evaluate_input_fn)  
model.predict(input_fn=predict_input_fn)
```

To the code!

kaggle.com/yufengg/census-with-w-d-estimators

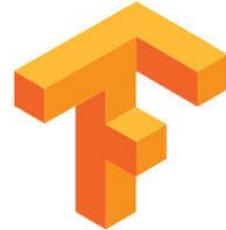




Typical structure

- Load data
- Set up feature columns
- Create your model
- Run the training loop (fit the model)
- Evaluate your model's accuracy (and other metrics)
- (optional) Predict new examples

Input data format



```
features = {
    'hours_per_week': array([16, 45, 50], dtype=int32),

    'relationship': SparseTensorValue(indices=array([[0, 0],[1, 0],[2, 0]]), values=array(['Not-in-family', ' Husband', ' Not-in-family']), dtype=object), shape=array([3, 1]),

    'sex': SparseTensorValue(indices=array([[0, 0],[1, 0],[2, 0]]), values=array([' Female', ' Male', ' Female']), dtype=object), shape=array([3, 1]),

    'age': array([49, 52, 31], dtype=int32)
    ...
}

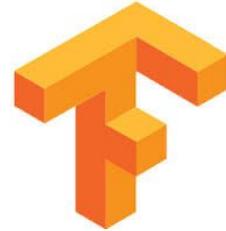
labels = [0 1 1]
```

Input data format



```
df = pandas.read_csv(filename)  
labels = df["labels"]  
  
return tf.estimator.inputs.pandas_input_fn(  
    x=df, y=labels)
```

Typical structure



- Load data
- **Set up feature columns**
- Create your model
- Run the training loop (fit the model)
- Evaluate your model's accuracy (and other metrics)
- (optional) Predict new examples

Feature columns



```
# Sparse base columns.  
  
sex = categorical_column_with_vocabulary_list(  
    key="sex", vocabulary_list=["female", "male"])  
  
education = categorical_column_with_hash_bucket(  
    key="education", hash_bucket_size=1000)  
  
...  
  
# Continuous base columns.  
  
age = numeric_column("age")  
  
...
```

Feature columns continued



```
# Transformations.  
  
age_buckets = bucketized_column(  
    age, boundaries=[18, 25, 30, 35, 40, 45, 50, 55, 60, 65])  
  
education_occupation = crossed_column(  
    ["education", "occupation"], hash_bucket_size=int(1e4))  
  
...  
  
# embeddings for deep learning  
  
embedding_column(workclass, dimension=8)
```

Feature columns continued



```
# Transformations.
```

```
age_buckets = layers.bucketized_column(  
    age, boundaries=[18, 25, 30, 35, 40, 45, 50, 55, 60, 65])
```

```
education_occupation = layers.crossed_column(  
    [education, occupation], hash_bucket_size=int(1e4))
```

```
...
```

```
# embedding  
layers.emt
```



```
ing  
kclas
```



Typical structure



- Load data
- Set up feature columns
- **Create your model**
- Run the training loop (fit the model)
- Evaluate your model's accuracy (and other metrics)
- (optional) Predict new examples



Make the model (Estimator)

```
m = tf.estimator.DNNLinearCombinedClassifier(  
    model_dir=model_dir,  
    linear_feature_columns=wide_columns,  
    dnn_feature_columns=deep_columns,  
    dnn_hidden_units=[100, 70, 50, 25])
```

Typical structure



- Load data
- Set up feature columns
- Create your model
- **Run the training loop**
- **Evaluate your model's accuracy (and other metrics)**
- (optional) Predict new examples



Train and Evaluate

```
m.train(input_fn=train_input_fn)  
results = m.evaluate(input_fn=eval_input_fn)  
  
print('Accuracy: %s' % results['accuracy'])
```

Make Predictions!



```
predict_input_fn =  
    tf.estimator.inputs.pandas_input_fn(  
        x=data_predict, batch_size=1,  
        num_epochs=1, shuffle=False)
```

```
predictions =  
    m.predict(input_fn=predict_input_fn)
```

Checkpointing and reloading a trained model

- Everything is stored in the `model_dir` folder

```
m = tf.estimator.DNNLinearCombinedClassifier(  
    model_dir=model_dir,  
    linear_feature_columns=wide_columns,  
    dnn_feature_columns=deep_columns,  
    dnn_hidden_units=[100, 70, 50, 25])
```

- If you run multiple `train` operations on the same `Estimator` and supply the same directory, training will resume where it left off.



Next: Tweak the TensorFlow model

Exploring the Wide & Deep Model's characteristics

Ideas: TensorFlow model tweaking

- Learning Rate
- Layers (# of layers, size of layers)
- Optimizer
- Feature crosses, buckets, embedding
- ... your dataset???

Building Custom Estimators for CNN models

Fashion-MNIST revisited

(Maybe CNNs will let us do better with
‘Fashion-MNIST’)

Brief intro to convolutional neural nets (CNNs)

Convolutional neural networks

- A fully-connected NN does not take into account spatial structure of the images
- What if instead we use a NN that tries to leverage spatial structure?

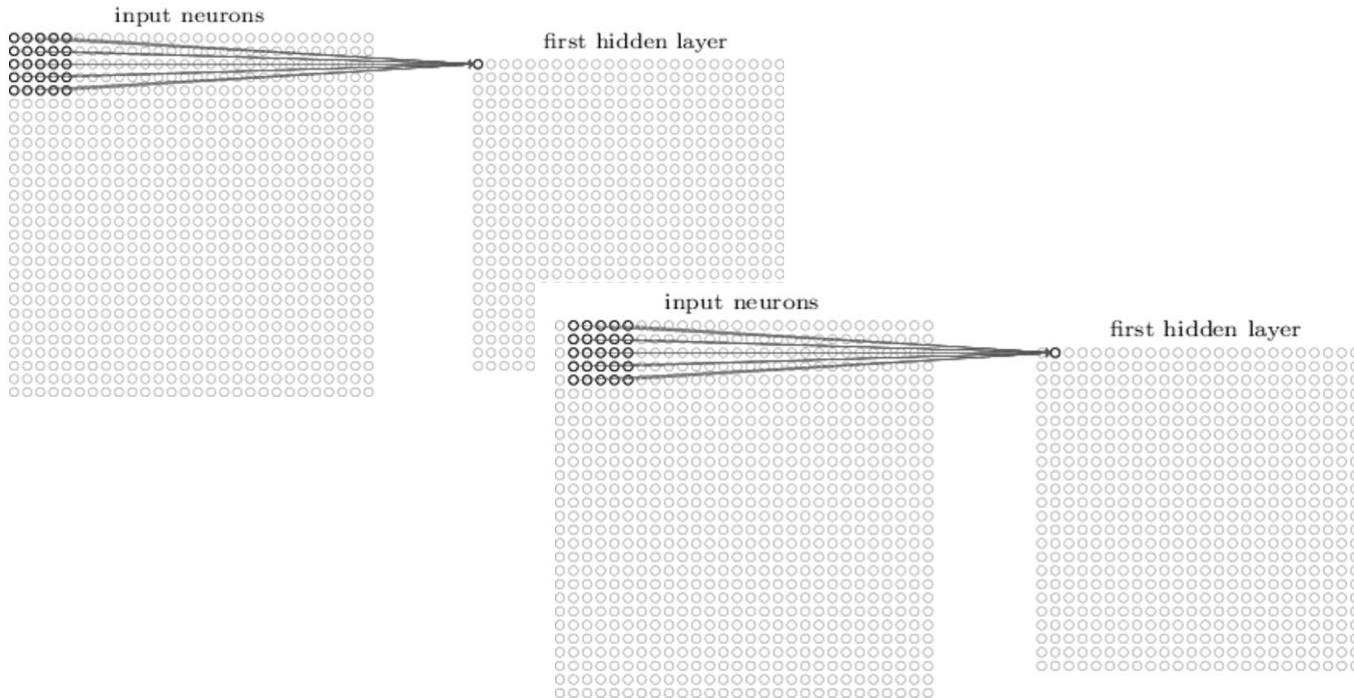


Convolutional neural networks

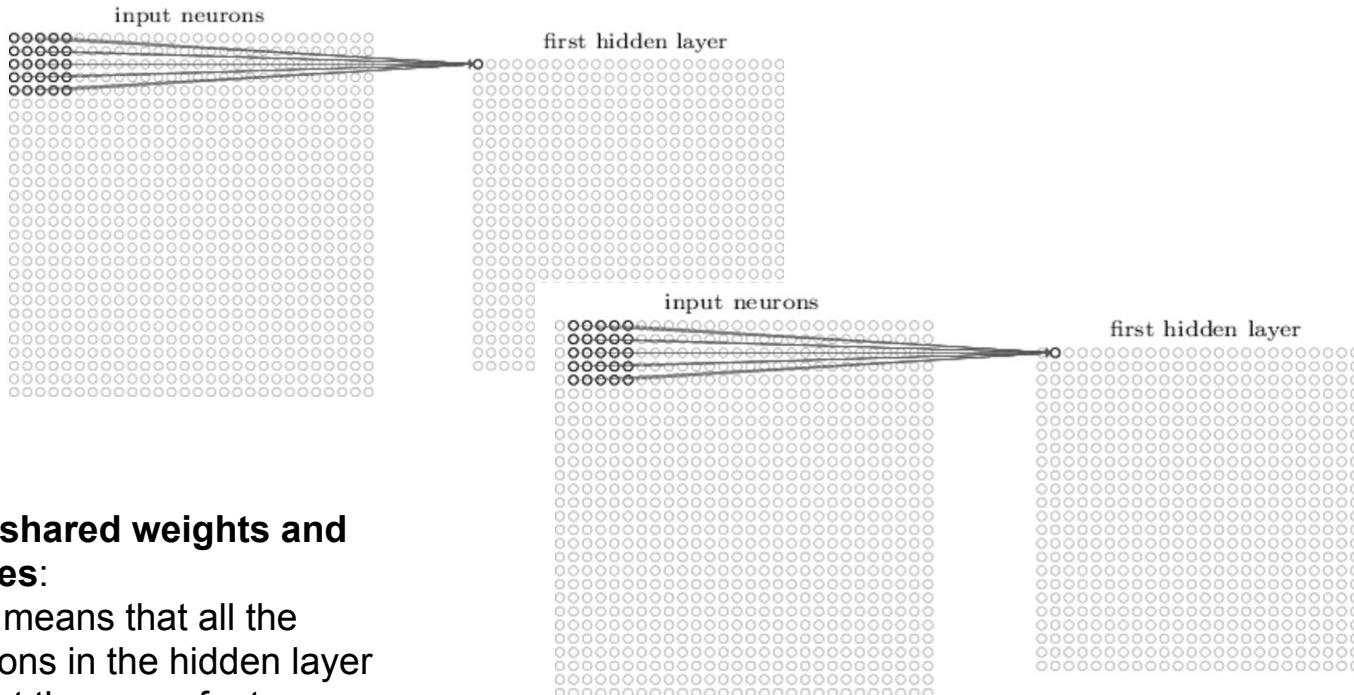
Three basic ideas:

- *local receptive fields*
- shared weights
- pooling





From: http://neuralnetworksanddeeplearning.com/chap6.html#introducing_convolutional_networks



Add **shared weights and biases**:

This means that all the neurons in the hidden layer detect the same feature

From: http://neuralnetworksanddeeplearning.com/chap6.html#introducing_convolutional_networks

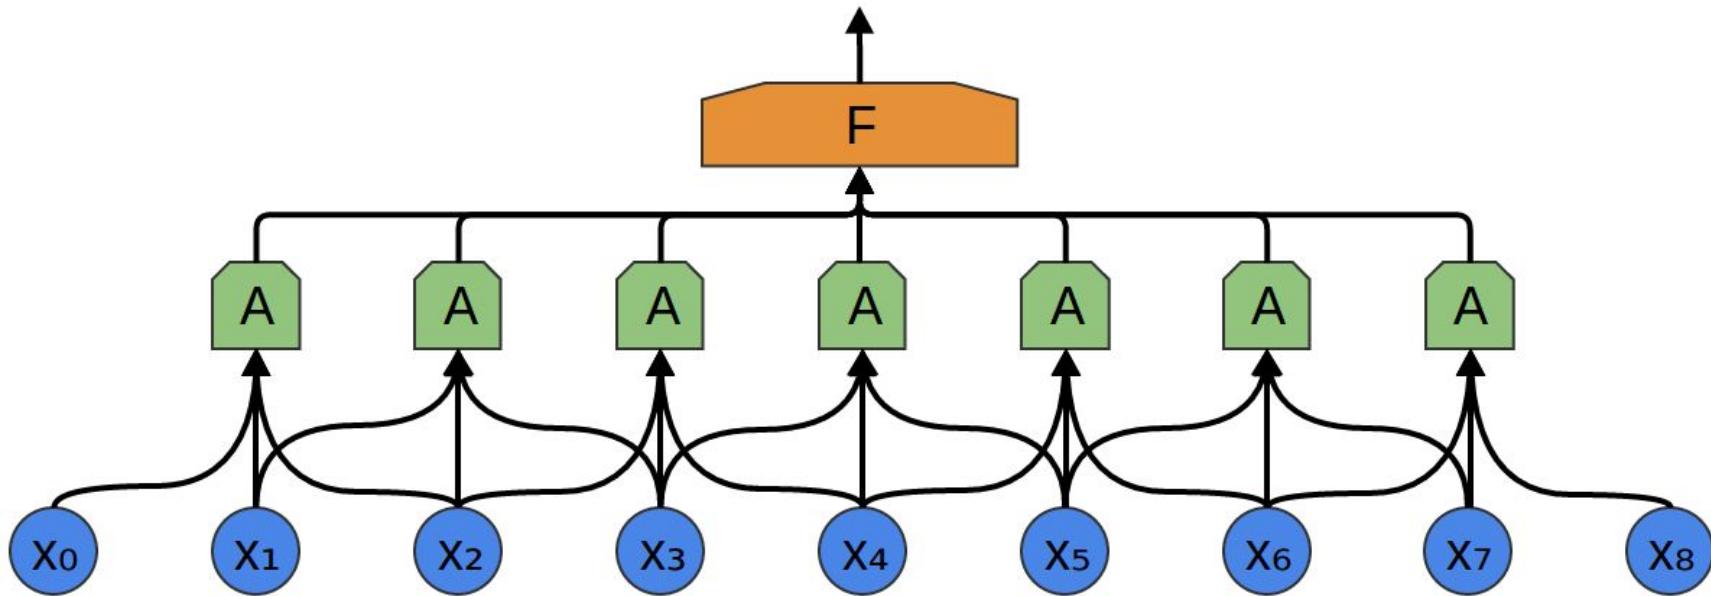


Image from: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>

1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0	0
0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1	0
0 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1	1
0	0	1	1	0
0	1	1	0	0

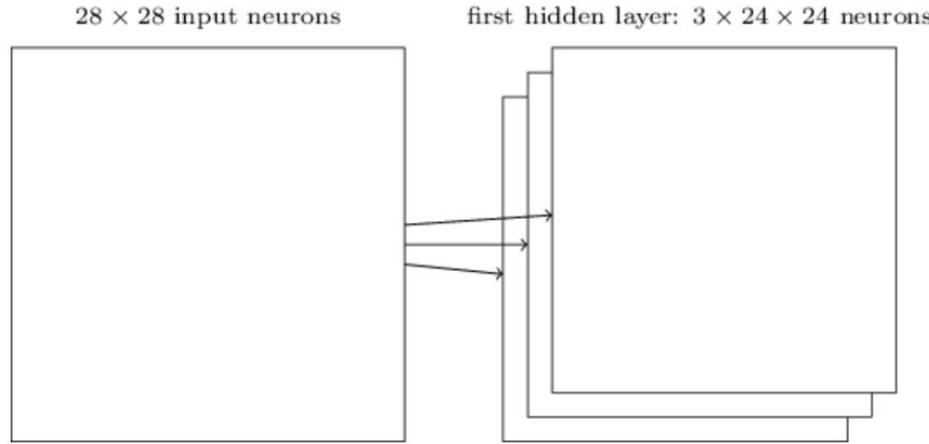
Image

4		

Convolved Feature

Convolution with 3×3 Filter. Source:

http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution



Typically a convolutional layer will have multiple different *feature maps*.

From: http://neuralnetworksanddeeplearning.com/chap6.html#introducing_convolutional_networks

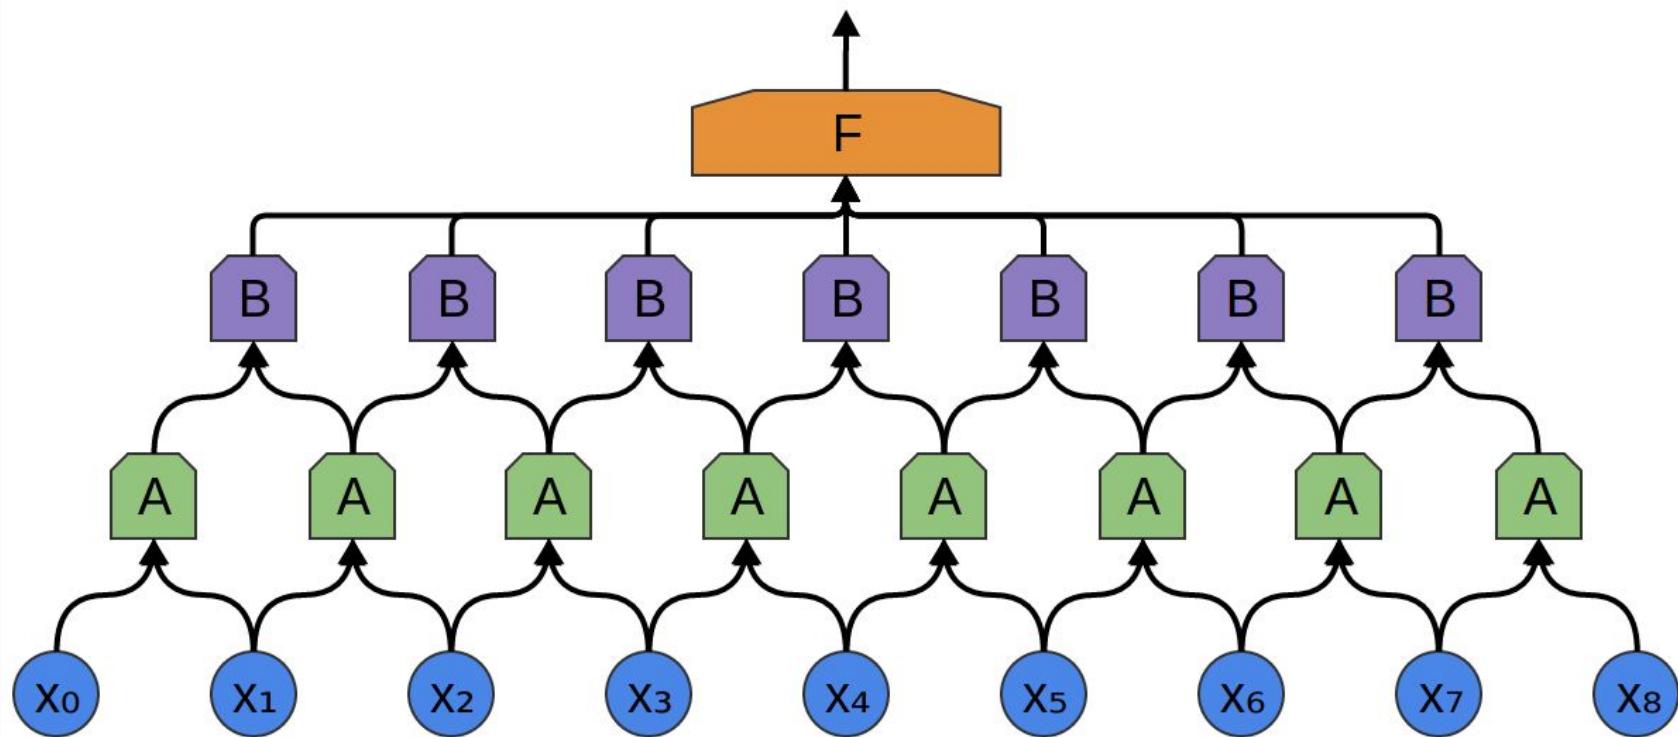


Image from: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>

Often
max-pooling
layers are
added

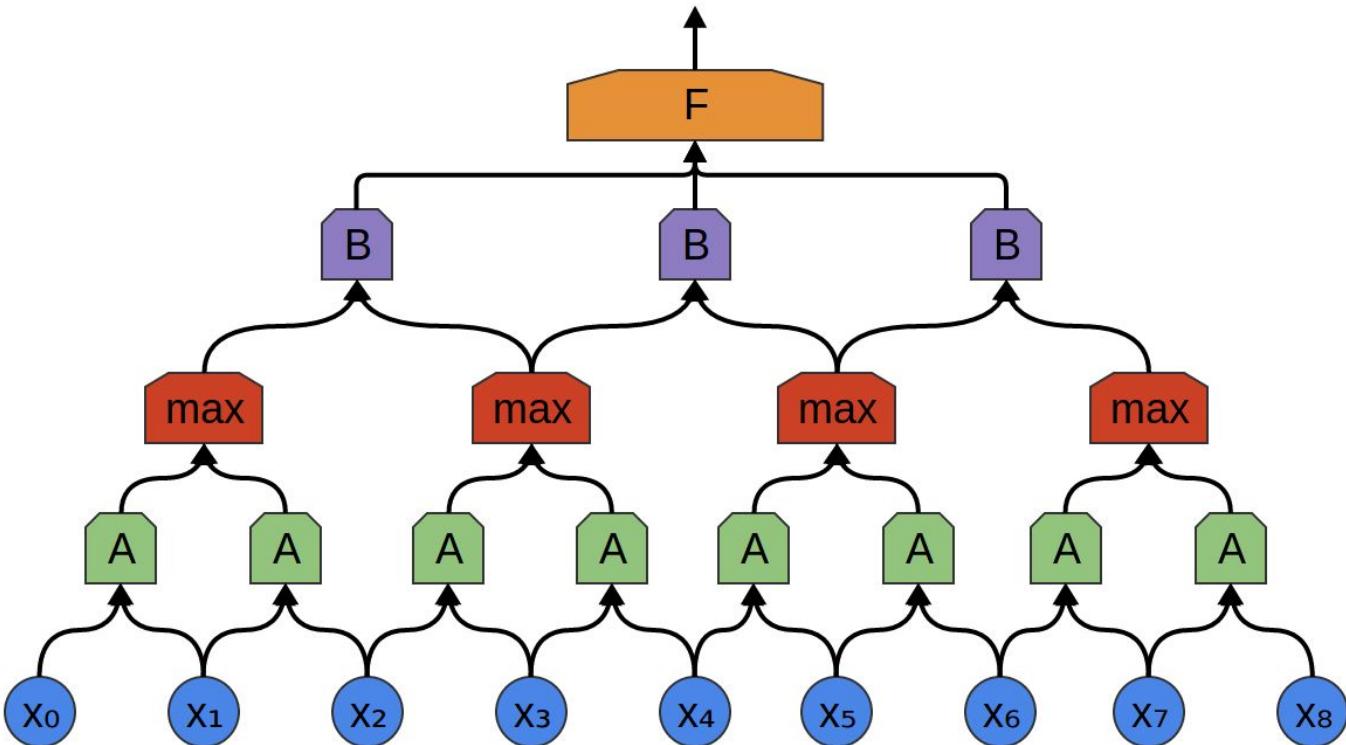
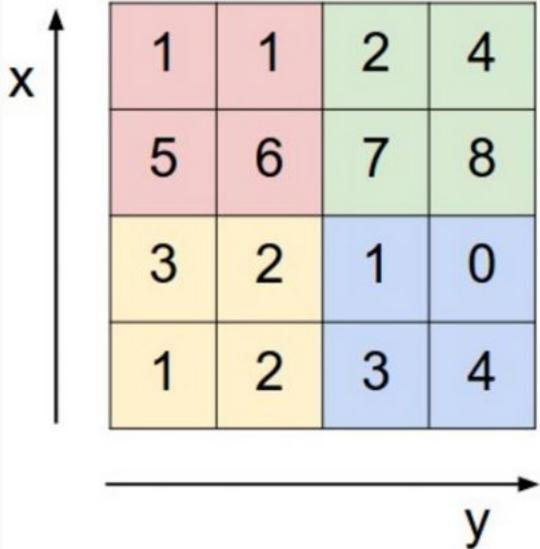


Image from: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>

Single depth slice



max pool with 2x2 filters
and stride 2

A 2x2 grid representing the output of max pooling. The values are:

6	8
3	4

Max pooling in CNN. Source: <http://cs231n.github.io/convolutional-networks/#pool>, via
<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

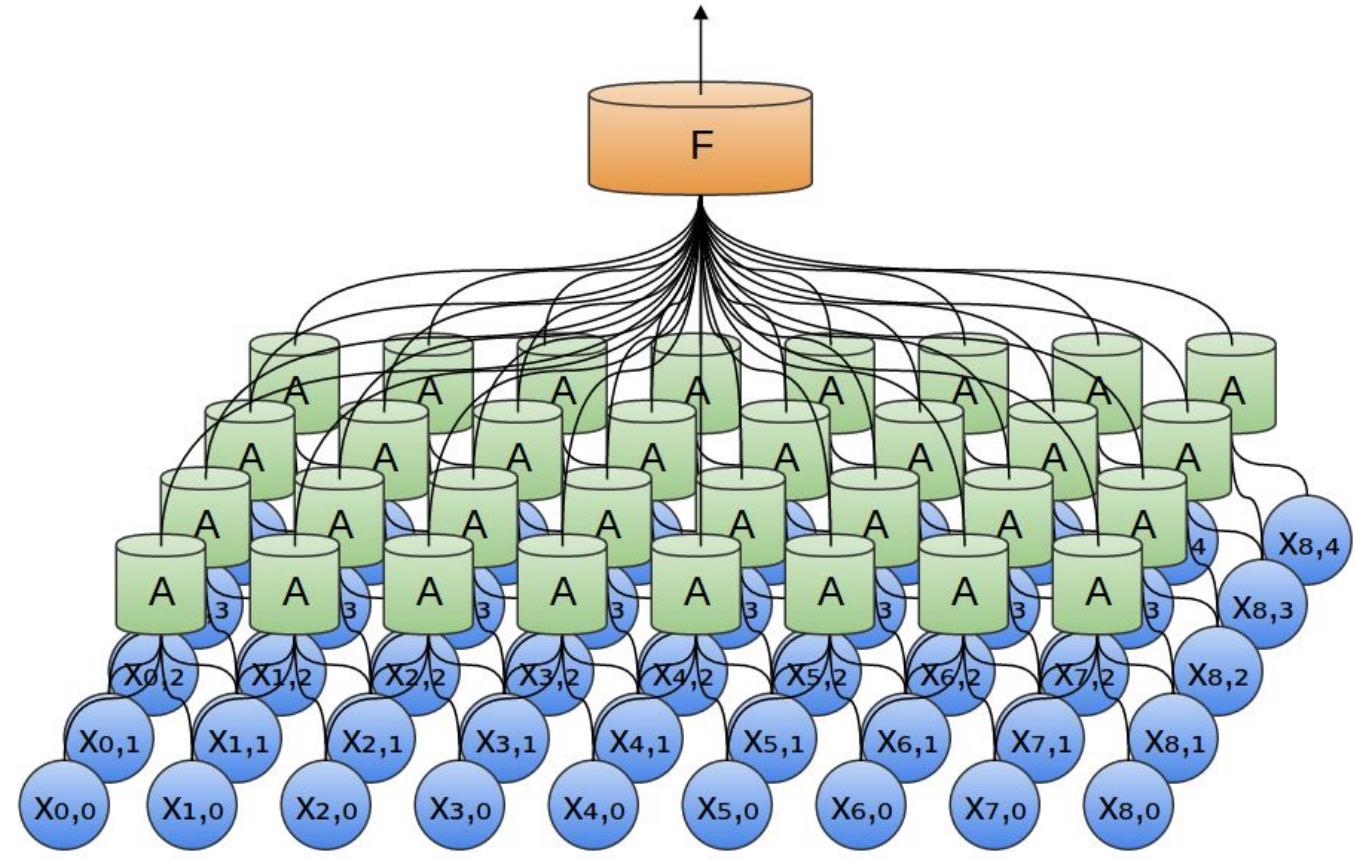


Image from: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>

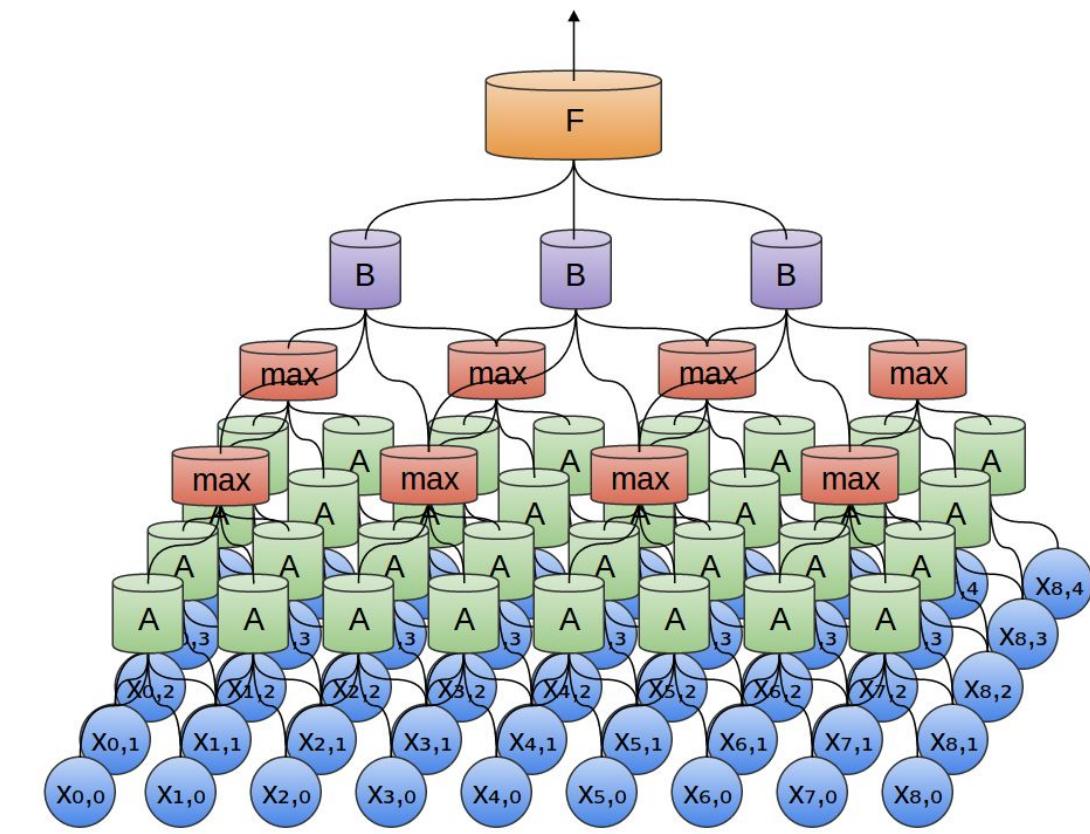
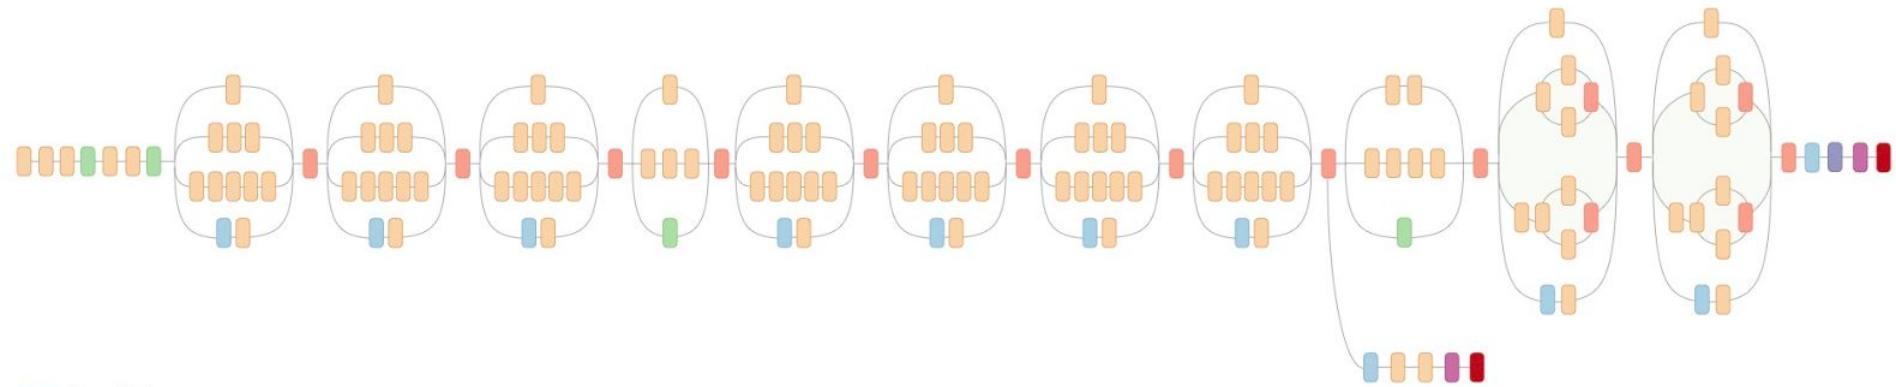


Image from: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>



- Convolution
- AvgPool
- MaxPool
- Concat
- Dropout
- Fully connected
- Softmax

From: http://googleresearch.blogspot.com/2016_03_01_archive.html

Let's build a custom **Estimator** for a CNN-based model

Lab: CNN Custom Estimators

Workshop section:

kaggle.com/yufengg/fashion-mnist-goes-cnn



These slides: bit.ly/tf-workshop-aicon

Create a Kaggle.com account

Go to kaggle.com/yufengg/kernels

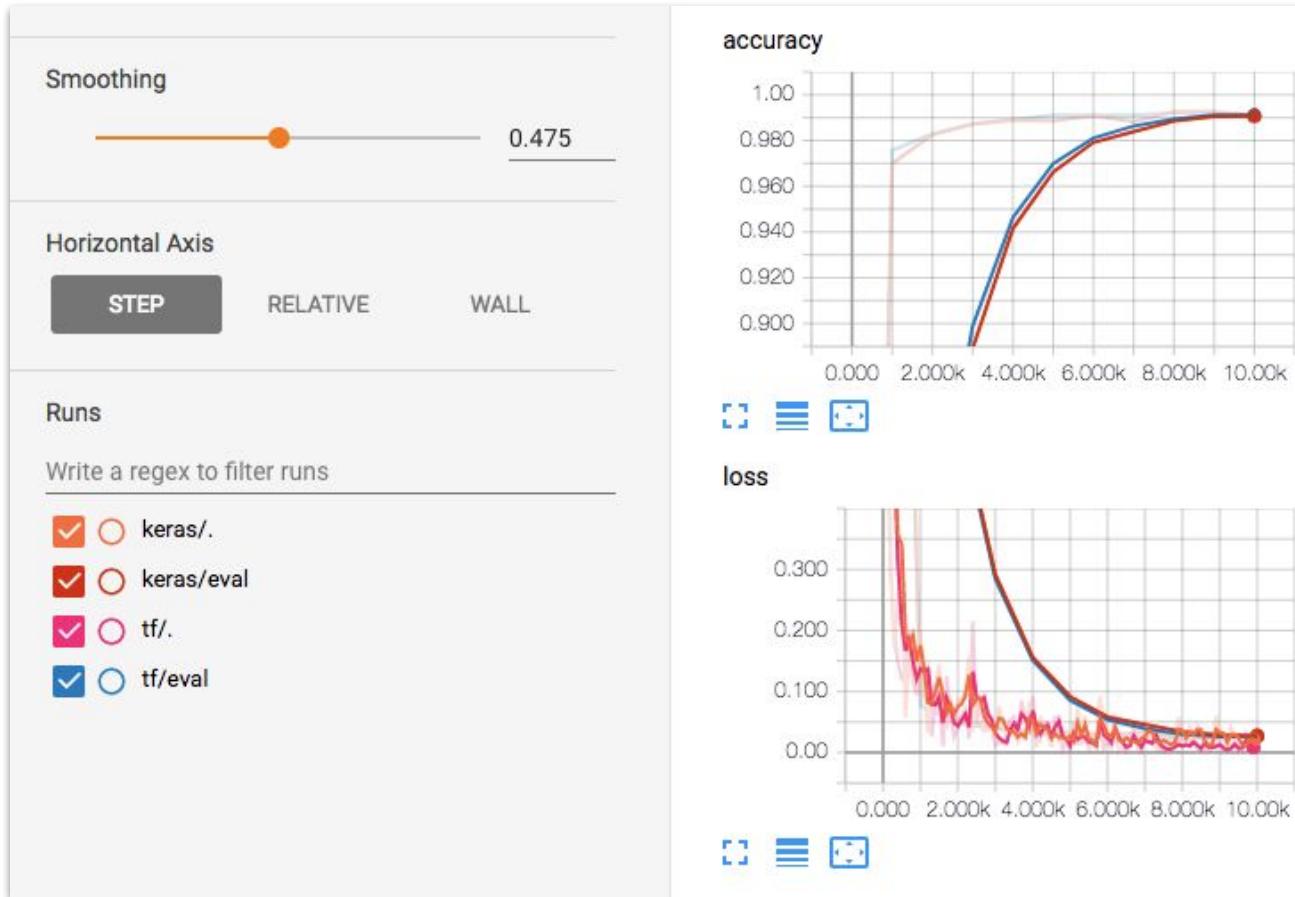
OR

Install TensorFlow & Jupyter: www.tensorflow.org/install/

pip install tensorflow pandas jupyter

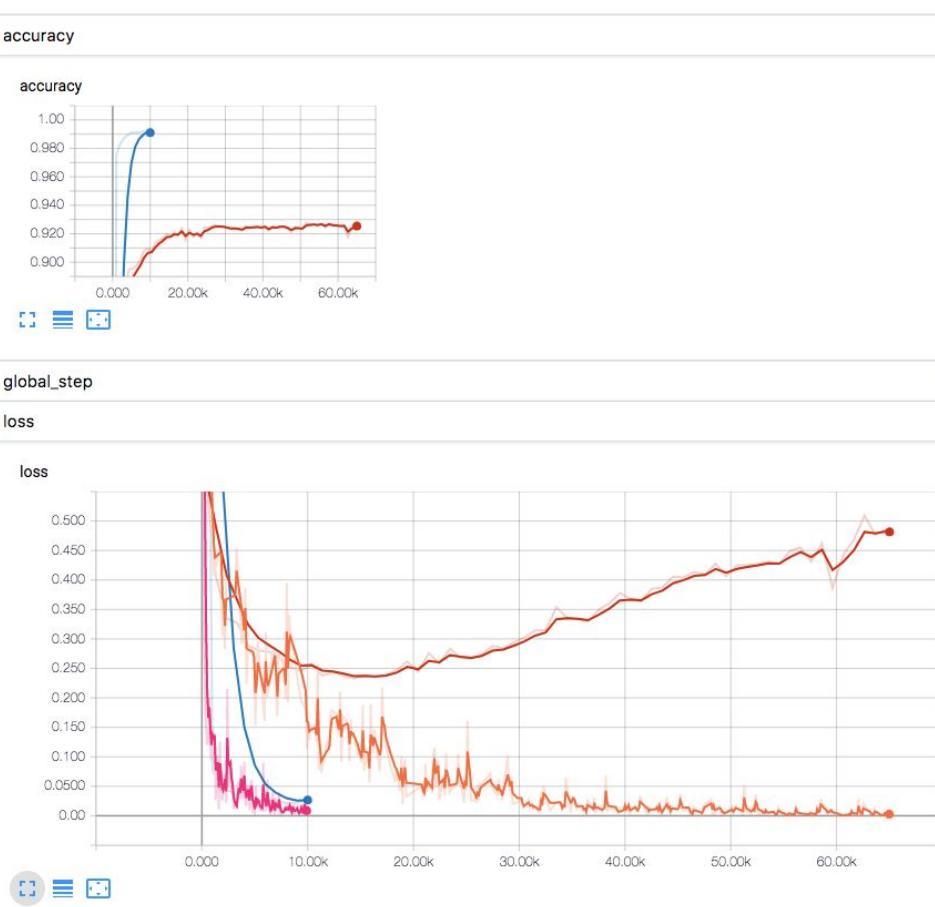
Clone GitHub Repo: bit.ly/tensorflow-workshop

Optional: download the ‘Fashion-MNIST’ files:
github.com/zalandoresearch/fashion-mnist#get-the-data



Write a regex to filter runs

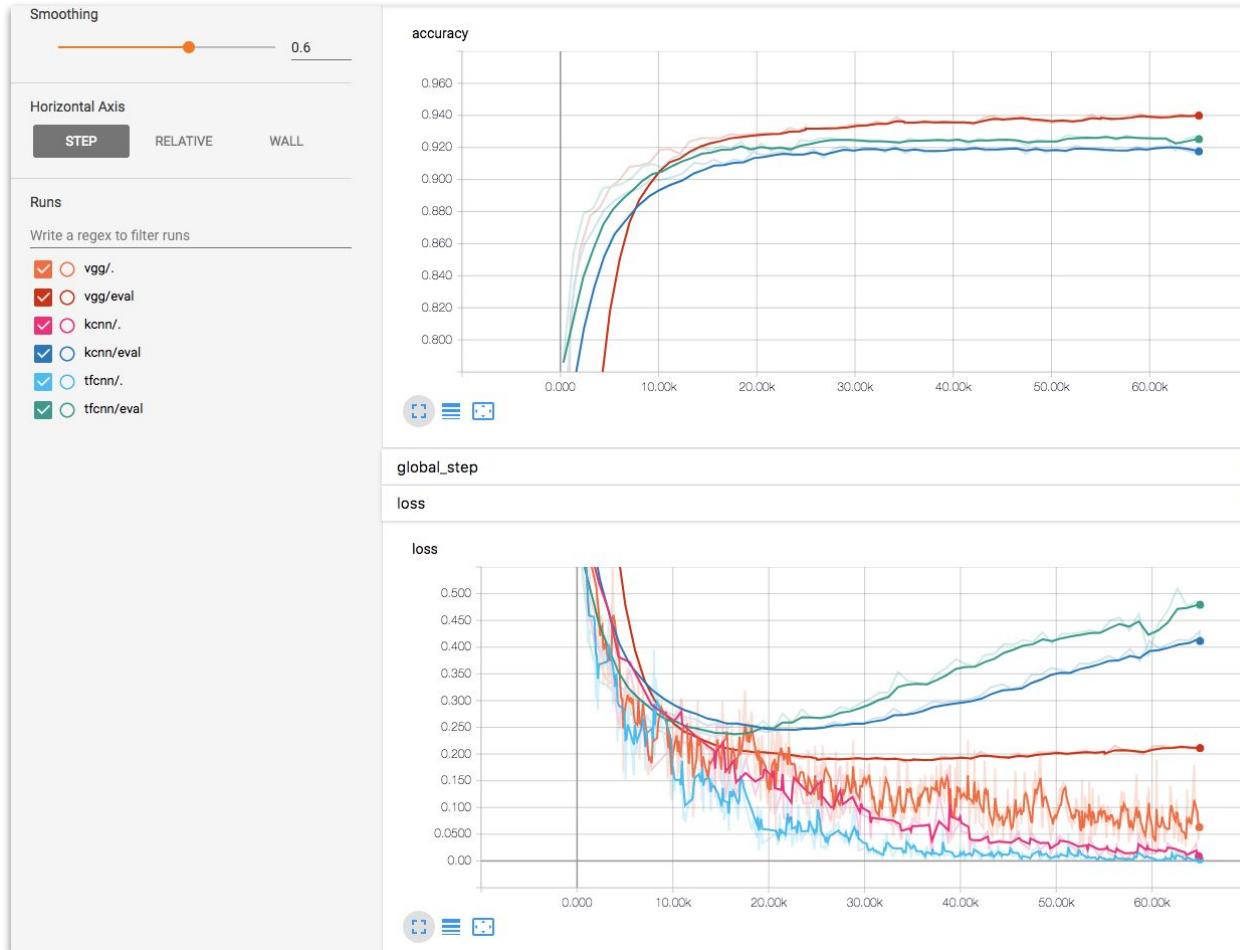
- fashion/.
- fashion/eval
- regmnist/.
- regmnist/eval



TOGGLE ALL RUNS

Comparing 'regular' with Fashion MNIST

Maybe there are better models?...



Wrapup



Google Cloud

AI Adventures

Presents:

What is Machine Learning?

Yufeng Guo

Developer Advocate



bit.ly/aiadventures

Exploring the art, science, and tools of machine learning

Thank you!

Rate this session:

bit.ly/tf-workshop-beij-ai18



Yufeng
yfg@google.com
[@YufengG](https://twitter.com/YufengG)
yufengg.com

What next?

Tutorials and code
tensorflow.org

TensorFlow Dev Summit
2017: <https://goo.gl/bv9mm7>
2018: tensorflow.org/dev-summit/

Intro to Deep Learning with TensorFlow
Udacity class goo.gl/iHssII

Stanford's CS231n
cs231n.github.io

TensorFlow Playground
goo.gl/mXhncM

Rate this session:
bit.ly/tf-workshop-beij-ai18



These Slides
bit.ly/tf-workshop-aicon

AI Adventures (on YouTube)
bit.ly/aiadventures



Deep Learning (Goodfellow, Bengio, Courville)
<http://www.deeplearningbook.org/>

Chris Olah's blog
colah.github.io

Michael Nielsen's book
neuralnetworksanddeeplearning.com

end